



# On multi-processor speed scaling with migration <sup>☆</sup>



Susanne Albers <sup>a,\*</sup>, Antonios Antoniadis <sup>b</sup>, Gero Greiner <sup>a</sup>

<sup>a</sup> Department of Computer Science, Technische Universität München, Boltzmannstr. 3, 85748 Garching, Germany

<sup>b</sup> Max-Planck-Institut für Informatik, Campus E1 4, 66123 Saarbrücken, Germany

## ARTICLE INFO

### Article history:

Received 8 March 2012

Accepted 31 January 2015

Available online 30 March 2015

### Keywords:

Energy efficiency

Offline algorithm

Online algorithm

Flow computation

Competitive analysis

## ABSTRACT

We investigate a very basic problem in dynamic speed scaling where a sequence of jobs, each specified by an arrival time, a deadline and a processing volume, has to be processed so as to minimize energy consumption. We study multi-processor environments with  $m$  parallel variable-speed processors assuming that job migration is allowed, i.e. whenever a job is preempted it may be moved to a different processor. We first study the offline problem and show that optimal schedules can be computed efficiently in polynomial time, given any convex non-decreasing power function. In contrast to a previously known strategy, our algorithm does not resort to linear programming. For the online problem, we extend two algorithms *Optimal Available* and *Average Rate* proposed by Yao et al. [15] for the single processor setting. Here we concentrate on power functions  $P(s) = s^\alpha$ , where  $s$  is the processor speed and  $\alpha > 1$  is a constant.

© 2015 Elsevier Inc. All rights reserved.

## 1. Introduction

Algorithmic techniques for energy savings in computing devices have received considerable research interest recently. A relatively new, effective approach is *dynamic speed scaling*. It relies on the fact that many modern microprocessors can operate at variable speed. Both Intel and AMD offer a range of such variable-speed processors. Obviously, high speed implies high performance. However, the higher the speed, the higher the energy consumption. The goal of dynamic speed scaling is to utilize the full speed/frequency spectrum of a processor and to use low speeds whenever possible.

In a seminal paper, initiating the algorithmic study of dynamic speed scaling, Yao, Demers and Shenker [15] introduced the following problem. Consider a sequence  $\sigma = J_1, \dots, J_n$  of  $n$  jobs that have to be scheduled on a variable-speed processor. Each job  $J_i$  is specified by a release time  $r_i$ , a deadline  $d_i$  and a processing volume  $w_i$ ,  $1 \leq i \leq n$ . In a feasible schedule,  $J_i$  must be processed within the time interval  $[r_i, d_i]$ . The processing volume  $w_i$  is the amount of work that must be finished to complete  $J_i$  and, intuitively, can be viewed as the number of required CPU cycles. The processing time of a job depends on the processor speed. If  $J_i$  is processed at speed  $s$ , then it takes  $w_i/s$  time units to complete the job. Preemption of jobs is allowed, i.e. the execution of a job may be stopped and resumed later. The well-known cube-root rule for CMOS devices states that the speed  $s$  of a processor is proportional to the cube-root of the power or, equivalently, that power is proportional to  $s^3$ . The algorithms literature considers generalizations of this rule. Early and in fact most of the previous

<sup>☆</sup> A preliminary version of this paper has appeared in the *Proceedings of the 23th Annual ACM Symposium on Parallelism in Algorithms and Architectures (SPAA)*, 279–288, 2011.

\* Corresponding author. Fax: +49 89 28917707.

E-mail addresses: [albers@in.tum.de](mailto:albers@in.tum.de) (S. Albers), [aantonio@mpi-inf.mpg.de](mailto:aantonio@mpi-inf.mpg.de) (A. Antoniadis), [greiner@informatik.tu-muenchen.de](mailto:greiner@informatik.tu-muenchen.de) (G. Greiner).

<sup>1</sup> Work supported by the German Research Foundation, AL 464/7-1.

work assumes that if a processor runs at speed  $s$ , then the required power is  $P(s) = s^\alpha$ , where  $\alpha > 1$  is a constant. More recent research even allows for general convex non-decreasing power functions  $P(s)$ . Obviously, energy is power integrated over time. The goal is to find a feasible schedule for the given job instance  $\sigma = J_1, \dots, J_n$  minimizing energy consumption.

The scheduling problem defined above has received the most research interest among algorithmic speed scaling problems. We will present the most important results below. Almost all of the previous studies assume that a single variable-speed processor is given. However, energy conservation and speed scaling techniques are equally interesting in multi-processor environments. Multi-core and many-core platforms will be the dominant processor architectures in the future. Nowadays many PCs and laptops are already equipped with dual-core and quad-core designs. Moreover, compute clusters and server farms, usually consisting of many high-speed processors, represent parallel multi-processor systems that have been used successfully in academia and enterprises for many years. Power dissipation has become a major concern in these environments. Finally, in research, multi-processor systems have always been investigated extensively, in particular as far as scheduling and resource management problems are concerned.

In this paper we study dynamic speed scaling in multi-processor environments. We adopt the framework by Yao et al. [15], as defined above, but assume that  $m$  parallel processors are given. Each processor can individually run at variable speed  $s$ ; the associated power function is  $P(s)$ . We consider the specific family of functions  $P(s) = s^\alpha$ , where  $\alpha > 1$ , as well as general convex non-decreasing functions  $P(s)$ . *Job migration* is allowed, i.e. whenever a job is preempted it may be moved to a different processor. Hence, over time, a job may be executed on various processors as long as the respective processing intervals do not overlap. Executing a job simultaneously on two or more processors is not allowed. The goal is to construct a feasible schedule minimizing the total energy consumption incurred by all the processors.

Two scenarios are of interest. In the offline setting, all jobs and their characteristics are known in advance. We wish to construct optimal schedules minimizing energy consumption. In the online setting, jobs arrive over time. Whenever a new job  $J_i$  arrives at time  $r_i$ , its deadline  $d_i$  and processing volume  $w_i$  are known. However, future jobs  $J_k$ , with  $k > i$ , are unknown. We use competitive analysis to evaluate the performance of online strategies [14]. An online algorithm  $A$  is called  $c$ -competitive if, for any job sequence  $\sigma$ , the energy consumption of  $A$  is at most  $c$  times the consumption of an optimal offline schedule.

**Previous work** We review results on dynamic speed scaling, focusing on deadline-based scheduling introduced by Yao et al. [15]. As mentioned above, most of the previous work addresses single-processor environments and power functions  $P(s) = s^\alpha$ , where  $\alpha > 1$ . In [15] Yao et al. first studied the offline problem and presented a polynomial time algorithm for computing optimal schedules. Refinements of the algorithm were given by Li et al. [11,12]. Yao et al. also presented two elegant online algorithms called *Optimal Available* and *Average Rate*. They proved that *Average Rate* achieves a competitive ratio of  $(2\alpha)^\alpha/2$ . The analysis is essentially tight as Bansal et al. [2] showed a nearly matching lower bound of  $((2-\delta)\alpha)^\alpha/2$ , where  $\delta$  goes to zero as  $\alpha$  tends to infinity. Bansal et al. [5] analyzed *Optimal Available* and, using a clever potential function, proved a competitiveness of exactly  $\alpha^\alpha$ . They also proposed a new strategy that attains a competitive ratio of  $2(\frac{\alpha}{\alpha-1})e^\alpha$ . As for lower bounds, Bansal et al. [4] showed that the competitiveness of any deterministic strategy is at least  $e^{\alpha-1}/\alpha$ . Irani et al. [9] proved that the offline algorithm by Yao et al. [15] also constructs optimal schedules if any convex non-decreasing power function  $P(s)$  is given. On the other hand, no online algorithm has been designed or analyzed for general convex power functions.

The framework by Yao et al. assumes that there is no upper bound on the allowed processor speed. Articles [3,7,13] study settings in which the processor has a maximum speed or only a finite set of discrete speed levels. Irani et al. [9] consider an extended problem where a variable-speed processor is equipped with an additional sleep state, assuming that even at speed zero a positive amount of energy is consumed.

The only previous work addressing deadline-based scheduling in multi-processor systems is [1,6,8,10]. Bingham and Greenstreet [6] show that, if job migration is allowed, the offline problem can be solved in polynomial time using linear programming. The result holds for general convex non-decreasing power functions. Lam et al. [10] study a setting with two speed-bounded processors. They show online algorithms that are constant competitive w.r.t. energy minimization and throughput maximization. Papers [1,8] assume that job migration is *not* allowed. In this case the offline problem is NP-hard, even if all jobs have the same processing volume [1]. A randomized  $B_\alpha$ -approximation algorithm and a randomized  $2(\frac{\alpha}{\alpha-1})e^\alpha B_\alpha$ -competitive online algorithm are given in [8]. Here  $B_\alpha$  is the  $\alpha$ -th Bell number. All the latter results were developed for the family of power functions  $P(s) = s^\alpha$ .

**Our contribution** In this paper we investigate dynamic speed scaling in general multi-processor environments, assuming that job migration is allowed. Using migration, scheduling algorithms can take advantage of the parallelism given by a multi-processor system in an effective way. We present a comprehensive study addressing the offline and, for the first time, also the online scenario.

First in Section 2 we study the offline problem and develop an efficient polynomial time algorithm for computing optimal schedules. The algorithm works for general convex non-decreasing power functions  $P(s)$ . As mentioned above, Bingham and Greenstreet [6] showed that the offline problem can be solved using linear programming. However, the authors mention that the complexity of their algorithm is too high for most practical applications. Instead in this paper we develop a strongly combinatorial algorithm that relies on repeated maximum flow computations. The approach might be helpful to solve other problems in scheduling and speed scaling.

Our algorithm is different from the single-processor strategy by Yao et al. [15]. In a series of phases, the algorithm partitions the jobs  $J_1, \dots, J_n$  into job sets  $\mathcal{J}_1, \dots, \mathcal{J}_p$  such that all jobs  $J_k \in \mathcal{J}_i$  are processed at the same uniform speed  $s_i$ ,  $1 \leq i \leq p$ . Each such job set is computed using maximum flow calculations. In order to construct a flow network, we have to identify various properties of a specific class of optimal schedules. A key property is that, knowing  $\mathcal{J}_1, \dots, \mathcal{J}_{i-1}$ , one can exactly determine the number of processors to be allocated to  $\mathcal{J}_i$ . At the beginning of the phase computing  $\mathcal{J}_i$ , the algorithm conjectures that the set  $\mathcal{J} = \{J_1, \dots, J_n\} \setminus (\mathcal{J}_1, \dots, \mathcal{J}_{i-1})$  of all remaining jobs forms the next set  $\mathcal{J}_i$ . If this turns out not to be the case, the algorithm repeatedly removes jobs  $J_k \in \mathcal{J}$  that do not belong to  $\mathcal{J}_i$ . The crucial step in the correctness proof is to show that each of these job removals is indeed correct so that, when the process terminates, the true set  $\mathcal{J}_i$  is computed.

In Section 3 we study the online problem and, as in the previous literature, focus on power functions  $P(s) = s^\alpha$ , where  $\alpha > 1$ . We adapt the two popular strategies *Optimal Available* and *Average Rate* to multi-processor environments. Algorithm *Optimal Available*, whenever a new job arrives, computes an optimal schedule for the remaining workload. This can be done using our new offline algorithm. We prove that *Optimal Available* is  $\alpha^\alpha$ -competitive, as in the single-processor setting. While the adaption of the algorithm is immediate, its competitive analysis becomes considerably more involved. We can extend the potential function analysis by Bansal et al. [5] but have to define a refined potential and have to prove several properties that specify how an optimal schedule changes in response to the arrival of a new job. As for the second algorithm *Average Rate*, we present a strategy that distributes load among the processors so that the densities  $\delta_i = w_i/(d_i - r_i)$  of the active jobs are almost optimally balanced. We prove that *Average Rate* achieves a competitiveness of  $(2\alpha)^\alpha/2 + 1$ . Hence, compared to competitive ratio in the single-processor setting, the factor only increases by the additive constant of 1.

## 2. A combinatorial offline algorithm

We develop a strongly combinatorial algorithm for constructing optimal offline schedules in polynomial time. Let  $\sigma = J_1, \dots, J_n$  be any job sequence and  $P(s)$  be an arbitrary convex non-decreasing power function. Lemma 1 below implies that there exist optimal schedules that use at most  $n$  different speeds, say speeds  $s_1 > s_2 > \dots > s_p$  where  $p \leq n$ . Our algorithm constructs such an optimal schedule in  $p$  phases, starting from an initially empty schedule  $S_0$ . Let  $S_{i-1}$  be the schedule obtained at the end of phase  $i - 1$ ,  $1 \leq i \leq p$ . In phase  $i$  the algorithm identifies the set  $\mathcal{J}_i$  of jobs that are processed at speed  $s_i$ ,  $1 \leq i \leq p$ . Schedule  $S_{i-1}$  is then extended by the jobs of  $\mathcal{J}_i$  to form a new schedule  $S_i$ . The job set  $\mathcal{J}_i$  and the extension of the schedule are determined using repeated maximum flow computations. Finally  $S_p$  is an optimal feasible schedule.

We present three lemmas that we will also use when analyzing an extension of *Optimal Available* in Section 3. In that section we will consider power functions  $P(s) = s^\alpha$ , where  $\alpha > 1$ .

**Lemma 1.** Any optimal schedule can be modified such that every job  $J_i$ ,  $1 \leq i \leq n$ , is processed at a constant non-varying speed. During the modification the schedule remains feasible and the energy consumption does not increase.

**Proof.** Consider any optimal schedule. While there exists a job  $J_i$ ,  $1 \leq i \leq n$ , that is processed at non-constant speed, let  $s_i$  be the average speed at which the job is executed. In the intervals in which  $J_i$  is processed modify the schedule so that  $J_i$  is executed at constant speed  $s_i$ . The total processing volume  $w_i$  of  $J_i$  is still completed in these intervals. By the convexity of the power function  $P(s)$  the modification does not increase the energy consumption of the schedule.  $\square$

By the above lemma we restrict ourselves to optimal schedules that process each job at a constant speed. Let  $S_{OPT}$  be the set of such optimal schedules. In any schedule  $S_{OPT} \in S_{OPT}$  the  $n$  jobs  $J_1, \dots, J_n$  can be partitioned into sets  $\mathcal{J}_1, \dots, \mathcal{J}_p$  such that all the jobs of  $\mathcal{J}_i$  are processed at the same speed  $s_i$ ,  $1 \leq i \leq p$ . Each job belongs to exactly one of these sets.

In the set  $S_{OPT}$  we identify a subset  $S'_{OPT}$  of schedules having some favorable properties. If the power function is  $P(s) = s^\alpha$ , with  $\alpha > 1$ , then let  $S'_{OPT} = S_{OPT}$ . Otherwise, if  $P(s)$  is a different power function, fix any  $\alpha > 1$  and let  $P_\alpha(s) = s^\alpha$  be an auxiliary power function. Let  $S'_{OPT} \subseteq S_{OPT}$  be the subset of the schedules that, among schedules of  $S_{OPT}$ , incur the smallest energy consumption if  $P(s)$  is replaced by  $P_\alpha(s)$ . That is, among the schedules in  $S_{OPT}$ , we consider those that would yield the smallest energy consumption if energy was accounted according to  $P_\alpha(s)$ . Since set  $S_{OPT}$  is non-empty, subset  $S'_{OPT}$  is also non-empty. Such an auxiliary power function was also used in [9] to break ties among optimal schedules.

The next lemma ensures that any schedule of  $S'_{OPT}$  can be modified such that the processor speeds change only at the release times and deadlines of jobs. Let  $\mathcal{I} = \{r_i, d_i \mid 1 \leq i \leq n\}$  be the set of all release times and deadlines. We consider the elements of  $\mathcal{I}$  in sorted order  $\tau_1 < \dots < \tau_{|\mathcal{I}|}$ , where  $|\mathcal{I}| \leq 2n$ . The time horizon in which jobs can be scheduled is  $[\tau_1, \tau_{|\mathcal{I}|})$ . We partition this time horizon along the job release times and deadlines into intervals  $I_j = [\tau_j, \tau_{j+1})$ ,  $1 \leq j < |\mathcal{I}|$ . Let  $|I_j| = \tau_{j+1} - \tau_j$  be the length of  $I_j$ .

**Lemma 2.** Given any optimal schedule  $S_{OPT} \in S'_{OPT}$ , in each interval  $I_j$  we can rearrange the schedule such that every processor uses a constant, non-varying speed in  $I_j$ ,  $1 \leq j < |\mathcal{I}|$ . During the modification the schedule remains feasible and the energy consumption with respect to  $P(s)$  and  $P_\alpha(s)$  does not increase.

**Proof.** Consider an arbitrary optimal schedule  $S_{OPT} \in \mathcal{S}'_{OPT}$  and let  $I_j$  be any interval,  $1 \leq j < |\mathcal{I}|$ . We show how to modify the schedule in  $I_j$ , without changing the energy consumption, so that the desired property holds for  $I_j$ . Modifying all the intervals in that way, we obtain the lemma.

For a given  $I_j$ , let  $S_{OPT}(I_j)$  be the schedule of  $S_{OPT}$  restricted to  $I_j$ . Moreover, let  $s_{j_1} > s_{j_2} > \dots > s_{j_l}$  be the different speeds employed in  $I_j$ . Assume that  $s_{j_k}$  is used for  $t_k$  time units in  $I_j$ , considering all the  $m$  processors, and let  $\mathcal{J}_{j_k}(I_j)$  be the set of jobs processed at speed  $s_{j_k}$  in  $I_j$ ,  $1 \leq k \leq l$ . Since jobs are processed at constant speed the sets  $\mathcal{J}_{j_k}(I_j)$ ,  $1 \leq k \leq l$ , are disjoint. For a job  $J_i \in \mathcal{J}_{j_k}(I_j)$ , let  $t_{i,j}$  be the total time for which  $J_i$  is scheduled in  $I_j$ . We have  $t_{i,j} \leq |I_j|$  since otherwise  $S_{OPT}(I_j)$  and hence  $S_{OPT}$  would not be feasible. Moreover,  $\sum_{J_i \in \mathcal{J}_{j_k}(I_j)} t_{i,j} = t_k$ .

We next construct a *working schedule*  $W$  of length  $\sum_{k=1}^l t_k$  for the jobs of  $\mathcal{J}_{j_1}(I_j), \dots, \mathcal{J}_{j_l}(I_j)$  and then distribute it among the processors. More specifically, for any  $k$  with  $1 \leq k \leq l$ , we construct a schedule  $W_k$  of length  $t_k$  in which the jobs of  $\mathcal{J}_{j_k}(I_j)$ , using a speed of  $s_{j_k}$ , are scheduled as follows: First the jobs  $J_i$  with  $t_{i,j} = |I_j|$  are processed, followed by the jobs  $J_i$  with  $t_{i,j} < |I_j|$ . Each  $J_i$  is assigned  $t_{i,j}$  consecutive time units in  $W_k$ . We concatenate the  $W_k$  to form the working schedule  $W = W_1 \circ \dots \circ W_l$  of length  $\sum_{k=1}^l t_k \leq m|I_j|$ . Next we distribute  $W$  among the  $m$  processors by assigning time window  $[(\mu-1)|I_j|, \mu|I_j|)$  of  $W$  to processor  $\mu$ , for  $1 \leq \mu \leq \lceil \sum_{k=1}^l t_k / |I_j| \rceil$ . Each window is scheduled from left to right on the corresponding processor, starting at the beginning of  $I_j$ . As we shall prove below, the total length of  $W$  is an integer multiple of  $|I_j|$  so that exactly  $\lceil \sum_{k=1}^l t_k / |I_j| \rceil$  processors are filled, without any idle period at the end. Let  $S'_{OPT}(I_j)$  be the new schedule for  $I_j$ . Furthermore, let  $S'_{OPT}$  be the schedule obtained from  $S_{OPT}$  when replacing  $S_{OPT}(I_j)$  by  $S'_{OPT}(I_j)$ . Schedule  $S_{OPT}$  is feasible because, for each job, the total execution time and employed speed have not changed. If in  $S'_{OPT}(I_j)$  the execution of a job  $J_i \in \mathcal{J}_{j_k}(I_j)$  is split among two processors  $\mu$  and  $\mu+1$ , then  $J_i$  is processed at the end of  $I_j$  on processor  $\mu$  and at the beginning of  $I_j$  on processor  $\mu+1$ . The two execution intervals do not overlap because  $t_{i,j} \leq |I_j|$ . Furthermore, the energy consumption of  $S'_{OPT}$  is the same as that of  $S_{OPT}$  because the processing intervals of the jobs have only been rearranged. This holds true with respect to both  $P(s)$  and  $P_\alpha(s)$ .

To establish the statement of the lemma for  $I_j$  we show that each  $t_k$  is an integer multiple of  $|I_j|$ , i.e.  $t_k = m_k |I_j|$  for some positive integer  $m_k$ ,  $1 \leq k \leq l$ . This implies that the new schedule never uses two or more speeds on any processor. So suppose that some  $t_k$  is not an integer multiple of  $|I_j|$ . Consider the smallest index  $k$  with this property. Then in  $S'_{OPT}(I_j)$ , on the corresponding processor, the processing of  $W_k$  ends at some time  $t$  strictly before the end of  $I_j$ , i.e.  $t < \tau_{j+1}$ . Furthermore, the job  $J_i$  handled last in  $W_k$  is processed for less than  $|I_j|$  time units since otherwise all jobs in  $W_k$  would be executed for  $|I_j|$  time units, contradicting the choice of  $k$ . So let  $t_{i,j} = |I_j| - \delta_i$ , for some  $\delta_i > 0$ . Set  $\delta = \min\{t - \tau_j, \tau_{j+1} - t, \delta_i, t_{i,j}\}$ . We now modify  $S'_{OPT}(I_j)$  on the processor handling the end of  $J_i$  within  $I_j$ . During the last  $\delta$  time units before  $t$  we decrease the speed by  $\epsilon$ , and during the first  $\delta$  time units after  $t$  we increase the speed by  $\epsilon$ , where  $\epsilon = (s_{j_k} - s_{j_{k+1}})/2$ . If  $k=l$ , we set  $s_{j_{k+1}} = 0$ , assuming that the processing of  $W_l$  is followed by an idle period. Since  $\delta \leq \tau_{j+1} - t$  and  $\delta \leq t - \tau_j$ , the modifications only affect the processor under consideration. Furthermore, the total work finished in  $[t - \delta, t + \delta]$  remains unchanged. As the speed was lowered in  $[t - \delta, t)$ , the first time units after  $t$  are used to finish the work on  $J_i$ . This can be accomplished before time  $t + \delta \leq t + \delta_i$  so that  $J_i$  is not processed for more than  $|I_j|$  time units and the resulting schedule is feasible.

Let  $S^m_{OPT}$  be the modified schedule obtained. Since  $P(s)$  is convex,  $S^m_{OPT}$  does not consume more energy than  $S'_{OPT}$  with respect to  $P(s)$ . Hence  $S^m_{OPT}$  is an optimal schedule. However, since  $\delta s_{j_k}^\alpha + \delta s_{j_{k+1}}^\alpha > 2\delta((s_{j_k} + s_{j_{k+1}})/2)^\alpha$  schedule  $S^m_{OPT}$  consumes strictly less energy than  $S'_{OPT}$  with respect to  $P_\alpha(s)$ . Finally, modify  $S^m_{OPT}$  so that each job is processed at constant speed. This can be done by processing job  $J_i$  and the job possibly scheduled after  $J_i$  in  $I_j$  using their respective average speeds. Since  $P(s)$  and  $P_\alpha(s)$  are convex, the respective energy consumptions do not increase. The resulting schedule belongs to set  $S_{OPT}$  and consumes strictly less energy than schedules  $S'_{OPT}$  and  $S_{OPT}$  with respect to  $P_\alpha(s)$ . We obtain a contradiction to the fact that  $S_{OPT} \in \mathcal{S}'_{OPT}$ .  $\square$

In the remainder of this section we consider optimal schedules  $S_{OPT} \in \mathcal{S}'_{OPT}$  satisfying the property of Lemma 2. Since set  $\mathcal{S}'_{OPT}$  is non-empty, these schedules always exist. As described above we will construct such a schedule  $S_{OPT}$  in phases, where phase  $i$  identifies the set  $\mathcal{J}_i$  of jobs processed at speed  $s_i$ ,  $1 \leq i \leq p$ . By Lemma 2, in each interval  $I_j$ ,  $1 \leq j < |\mathcal{I}|$ , the sets  $\mathcal{J}_1, \dots, \mathcal{J}_p$  occupy different processors, i.e. there is no processor handling jobs of two different sets. Hence, in phase  $i$ , when determining  $\mathcal{J}_i$  and extending the previous schedule  $S_{i-1}$ , we only need to know the number of processors occupied by  $\mathcal{J}_1, \dots, \mathcal{J}_{i-1}$  in any interval  $I_j$ ,  $1 \leq j < |\mathcal{I}|$ . The exact assignment of the corresponding jobs to the reserved processors is irrelevant though, as  $\mathcal{J}_i$  does not share processors with the previous sets. In determining  $\mathcal{J}_i$  a crucial question is, how many processors to allocate to the jobs of  $\mathcal{J}_i$  in any interval  $I_j$ ,  $1 \leq j < |\mathcal{I}|$ . Here the following lemma is essential.

A job  $J_k$ ,  $1 \leq k \leq n$ , is called *active* in  $I_j$  if  $I_j \subseteq [r_k, d_k]$ , i.e. the time period in which  $J_k$  can be scheduled includes  $I_j$ . Given schedule  $S_{OPT}$ , let  $n_{ij}$  be the number of jobs of  $\mathcal{J}_i$  that are active in  $I_j$ . Furthermore, let  $m_{ij}$  denote the number of processors occupied by  $\mathcal{J}_i$  in  $I_j$ .

**Lemma 3.** Let  $S_{OPT} \in \mathcal{S}'_{OPT}$  be any schedule satisfying the property of Lemma 2. Then the jobs of  $\mathcal{J}_i$  occupy  $m_{ij} = \min\{n_{ij}, m - \sum_{l=1}^{i-1} m_{lj}\}$  processors in  $I_j$ , where  $1 \leq i \leq p$  and  $1 \leq j < |\mathcal{I}|$ .

**Proof.** Consider a set  $\mathcal{J}_i$  and an interval  $I_j$ . Obviously, the  $n_{ij}$  jobs of  $\mathcal{J}_i$  that are active in  $I_j$  cannot occupy more than  $n_{ij}$  processors using a positive speed  $s_i$  throughout  $I_j$ . Furthermore, the jobs can only occupy the  $m - \sum_{l=1}^{i-1} m_{lj}$  processors not taken by  $\mathcal{J}_1, \dots, \mathcal{J}_{i-1}$ . We show that in fact  $m_{ij} = \min\{n_{ij}, m - \sum_{l=1}^{i-1} m_{lj}\}$  processors are used by  $\mathcal{J}_i$ .

So suppose that  $m'_{ij}$  processors, where  $m'_{ij} < m_{ij}$ , are used. Then, since  $m'_{ij} < m - \sum_{l=1}^{i-1} m_{lj}$ , there exists at least one processor  $P$  running at speed  $s$  with  $s < s_i$  in  $I_j$ . Consider the schedule at the beginning of  $I_j$ . Choose a  $\delta, \delta > 0$ , such that the  $m'_{ij}$  processors handling  $\mathcal{J}_i$  and processor  $P$  do not preempt jobs in  $[\tau_j, \tau_j + \delta)$ , i.e. they each handle at most one job in that time window. In fact  $P$  might not handle a job if  $s = 0$ . As  $m'_{ij} < n_{ij}$ , there must exist one job  $J_k \in \mathcal{J}_i$  that is active in  $I_j$  but not scheduled within  $[\tau_j, \tau_j + \delta)$  on any of the  $m'_{ij}$  processors using speed  $s_i$ . This job is not scheduled on any other processor within  $[\tau_j, \tau_j + \delta)$  either, because the other processors run at speeds higher or lower than  $s_i$ . Thus  $J_k$  is not executed on any processor within  $[\tau_j, \tau_j + \delta)$ . In the entire schedule, consider a processor and an associated time window  $W$  of length  $\delta' \leq \delta$  handling  $J_k$ . We now modify the schedule by reducing the speed in  $W$  by  $\epsilon$ , where  $\epsilon = (s_i - s)/2$ . At the same time we increase the speed on processor  $P$  in  $[\tau_j, \tau_j + \delta')$  by  $\epsilon$  and use the extra processing capacity to complete the missing portion of  $J_k$  not finished in  $W$ . Let  $S_{OPT}^m$  be the modified schedule. Since  $P(s)$  is convex,  $S_{OPT}^m$  does not consume more energy than  $S_{OPT}$  with respect to  $P(s)$ . With respect to  $P_\alpha(s)$ , schedule  $S_{OPT}^m$  consumes less energy than  $S_{OPT}$  because  $\delta' s_i^\alpha + \delta' s^\alpha > 2\delta'((s_i + s)/2)^\alpha$ . Finally, modify  $S_{OPT}^m$  so that each job is processed at a constant speed. We obtain a schedule that belongs to set  $S_{OPT}$  and consumes strictly less energy than schedule  $S_{OPT}$  with respect to  $P_\alpha(s)$ . This contradicts the fact that  $S_{OPT} \in S'_{OPT}$ .  $\square$

**Lemma 3** has an interesting implication: Suppose that job sets  $\mathcal{J}_1, \dots, \mathcal{J}_{i-1}$  along with the number of occupied processors  $m_{1j}, \dots, m_{i-1,j}$ , for  $1 \leq j < |\mathcal{I}|$ , have been determined. Furthermore, suppose that the set  $\mathcal{J}_i$  is known. Then, using **Lemma 3**, we can immediately determine the number  $m_{ij}$  of processors used by  $\mathcal{J}_i$  in  $I_j$ ,  $1 \leq j < |\mathcal{I}|$ . Moreover, we can compute the speed  $s_i$  by observing that, since  $P(s)$  is non-decreasing,  $s_i$  can be set to the minimum average speed necessary to complete the jobs of  $\mathcal{J}_i$  in the reserved processing intervals. More precisely, let  $W_i = \sum_{J_k \in \mathcal{J}_i} w_k$  be the total processing volume of  $\mathcal{J}_i$  and  $P_i = \sum_{1 \leq j < |\mathcal{I}|} m_{ij}|I_j|$  be the total length of the reserved processing intervals. Then  $s_i = W_i/P_i$ . We will make use of this fact when identifying  $\mathcal{J}_i$ .

**Description of the algorithm** We show how each phase  $i$ , that determines job set  $\mathcal{J}_i$ , works. Assume again that  $\mathcal{J}_1, \dots, \mathcal{J}_{i-1}$  along with processor numbers  $m_{1j}, \dots, m_{i-1,j}$ , for  $1 \leq j < |\mathcal{I}|$ , are given. In phase  $i$  the algorithm operates in a series of rounds, maintaining a job set  $\mathcal{J}$  that represents the current estimate for the true  $\mathcal{J}_i$ . At any time the invariant  $\mathcal{J}_i \subseteq \mathcal{J}$  holds. Initially, prior to the first round, the algorithm sets  $\mathcal{J} := \{J_1, \dots, J_n\} \setminus (\mathcal{J}_1 \cup \dots \cup \mathcal{J}_{i-1})$ , which is the set of all remaining jobs, and the invariant is obviously satisfied. In each round the algorithm checks if the current  $\mathcal{J}$  is the desired  $\mathcal{J}_i$ . This can be verified using a maximum flow computation. If  $\mathcal{J}$  turns out to be  $\mathcal{J}_i$ , then phase  $i$  stops. Otherwise the algorithm determines a job  $J_k \in \mathcal{J} \setminus \mathcal{J}_i$ , removes that job from  $\mathcal{J}$  and starts the next round in which the updated set  $\mathcal{J}$  is checked.

In the following we describe the maximum flow computation invoked for a given  $\mathcal{J}$ . First the algorithm determines the number of processors to be allocated to  $\mathcal{J}$  in each interval. For any  $I_j$ ,  $1 \leq j < |\mathcal{I}|$ , let  $n_j$  be the number of jobs in  $\mathcal{J}$  that are active in  $I_j$ . According to **Lemma 3** the algorithm reserves  $m_j = \min\{n_j, m - \sum_{l=1}^{i-1} m_{lj}\}$  processors in  $I_j$ . These numbers form a vector  $\vec{m} = (m_j)_{1 \leq j < |\mathcal{I}|}$ . The speed is set to  $s = W/P$ , where  $W = \sum_{J_k \in \mathcal{J}} w_k$  is the total processing volume and  $P = \sum_{1 \leq j < |\mathcal{I}|} m_j|I_j|$  is the reserved processing time.

Next we define the graph  $G(\mathcal{J}, \vec{m}, s)$  of the maximum flow computation. For each job  $J_k \in \mathcal{J}$  we introduce a vertex  $u_k$ , and for each interval  $I_j$  with  $m_j > 0$  we introduce a vertex  $v_j$ . For any  $u_k$  and  $v_j$  such that  $J_k$  is active in  $I_j$  we add a directed edge  $(u_k, v_j)$  of capacity  $|I_j|$ . Hence each  $u_k$  is connected to exactly those  $v_j$  such that  $J_k$  can be scheduled in  $I_j$ . The edge capacity of  $(u_k, v_j)$  is equal to  $|I_j|$ , i.e. the maximum time for which a job can be processed in  $I_j$ . Furthermore, we introduce a source vertex  $u_0$  that is connected to every  $u_k$  on a directed edge  $(u_0, u_k)$ . The edge capacity is set to  $w_k/s$ , which is the total processing time needed for  $J_k$  using speed  $s$ . Finally we introduce a sink  $v_0$  and connect each  $v_j$  to  $v_0$  via a directed edge  $(v_j, v_0)$  of capacity  $m_j|I_j|$ , which is the total processing time available on the reserved processors in  $I_j$ . The global structure of  $G(\mathcal{J}, \vec{m}, s)$  is depicted in **Fig. 1**. The value of a maximum flow in  $G(\mathcal{J}, \vec{m}, s)$  is upper bounded by  $F_G = \sum_{J_k \in \mathcal{J}} w_k/s$  because this is the total capacity of the edges leaving the source  $u_0$ . This is equal to the total capacity of the edges leading into the sink  $v_0$  because that value is

$$\sum_{1 \leq j < |\mathcal{I}|} m_j|I_j| = P = W/s = \sum_{J_k \in \mathcal{J}} w_k/s = F_G.$$

We argue that job set  $\mathcal{J}$  using a constant speed of  $s$  can be feasibly scheduled on the reserved processors represented by  $\vec{m}$  if and only if there exists a maximum flow of value  $F_G$  in  $G(\mathcal{J}, \vec{m}, s)$ . First suppose that there exists a feasible schedule using speed  $s$ . Then each  $J_k \in \mathcal{J}$  is processed for  $w_k/s$  time units and we send  $w_k/s$  units of flow along edge  $(u_0, u_k)$ . The  $w_k/s$  processing units of  $J_k$  are scheduled in intervals  $I_j$  in which the job is active. If  $J_k$  is processed for  $t_{kj}$  time units in  $I_j$ , then we send  $t_{kj}$  units of flow along edge  $(u_k, v_j)$ . The edge capacity is observed because  $t_{kj} \leq |I_j|$ . The total amount of flow entering and hence leaving any  $v_j$  is equal to the total processing time on the  $m_j$  reserved processors in  $I_j$ . This value is equal to  $m_j|I_j|$  because, by the choice of  $s$ , all reserved processors are busy throughout the execution intervals. Hence the amount of flow sent along  $(v_j, v_0)$  is equal to the edge capacity.



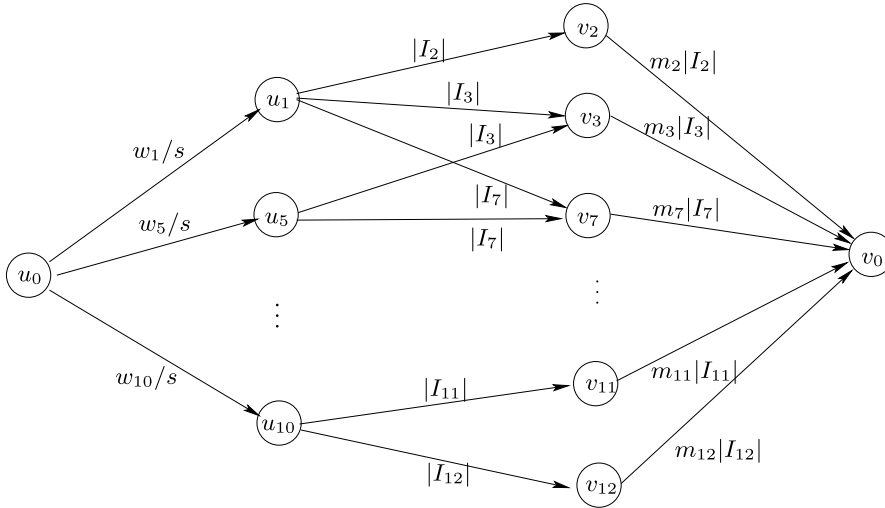


Fig. 1. The basic structure of  $G(\mathcal{J}, \bar{m}, s)$ , assuming that set  $\mathcal{J} = \{J_1, J_5, \dots, J_{10}\}$  can be scheduled in intervals  $I_2, I_3, I_7, \dots, I_{11}, I_{12}$ .

#### Algorithm Optimal Schedule

1.  $\mathcal{J} := \{J_1, \dots, J_n\}$ ;  $i := 0$ ;  $S_0 :=$  empty schedule;
2. **while**  $\mathcal{J} \neq \emptyset$  **do**
3.    $i := i + 1$ ; set\_found := false;
4.   **while**  $\neg$  set\_found **do**
5.      $n_j :=$  number of jobs of  $\mathcal{J}$  that are active in  $I_j$ , for  $1 \leq j < |\mathcal{I}|$ ;
6.      $m_j := \min\{n_j, \sum_{l=1}^{i-1} m_{lj}\}$ ;  $\bar{m} := (m_j)_{1 \leq j < |\mathcal{I}|}$ ;
7.      $W := \sum_{J_k \in \mathcal{J}} w_k$ ;  $P := \sum_{1 \leq j < |\mathcal{I}|} m_j |I_j|$ ;  $s := W/P$ ;
8.     Compute the value  $F$  of a maximum flow in  $G(\mathcal{J}, \bar{m}, s)$ ;
9.     **if**  $F = W/s$  **then** set\_found := true;
10.    **else** Determine edge  $(v_j, v_0)$  carrying less than  $m_j |I_j|$  units of flow and edge  $(u_k, v_j)$  carrying less than  $|I_j|$  units of flow; Set  $\mathcal{J} := \mathcal{J} \setminus \{J_k\}$ ;
11.     $\mathcal{J}_i := \mathcal{J}$ ;  $s_i := s$ ;  $m_{ij} := m_j$  for  $1 \leq j < |\mathcal{I}|$ ;
12.    Extend  $S_{i-1}$  by a feasible schedule for  $\mathcal{J}_i$  using speed  $s_i$  and  $m_{ij}$  processors in  $I_j$ ;
13.    Let  $S_i$  be the resulting schedule;
13.     $\mathcal{J} := \{J_1, \dots, J_n\} \setminus (\mathcal{J}_1 \cup \dots \cup \mathcal{J}_i)$ ;

Fig. 2. The entire offline algorithm for computing an optimal schedule.

On the other hand, assume that a flow of value  $F_G$  is given. If  $t_{kj}$  units of flow are routed along  $(u_k, v_j)$ , we process  $J_k$  for  $t_{kj}$  time units in  $I_j$ . This is possible because  $J_k$  is active in  $I_j$  and  $t_{kj} \leq |I_j|$ . The capacity constraints on the edge  $(v_j, v_0)$  ensure that not more than  $m_j |I_j|$  time units are assigned to  $I_j$ . Moreover, since the flow value is  $F_G$ , each edge  $(u_0, u_k)$  is saturated and hence  $J_k$  is fully processed for  $w_k/s$  time units at speed  $s$ . Within each  $I_j$  we can easily construct a feasible schedule by concatenating the processing intervals of length  $t_{kj}$  associated with the jobs  $J_k$  active in  $I_j$ . The resulting sequential schedule, say  $S$ , of length  $\sum_{J_k \in \mathcal{J}} t_{kj} = m_j |I_j|$  is split among the  $m_j$  reserved processors by assigning time range  $[(\mu - 1)|I_j|, \mu |I_j|)$  of  $S$  to the  $\mu$ -th reserved processor,  $1 \leq \mu \leq m_j$ .

We proceed to describe how the algorithm determines  $\mathcal{J}_i$ . In each round, for the current set  $\mathcal{J}$ , the algorithm invokes a maximum flow computation in the graph  $G(\mathcal{J}, \bar{m}, s)$  defined in the previous paragraphs. If the flow value is equal to  $F_G$ , then the algorithm stops and sets  $\mathcal{J}_i := \mathcal{J}$ . We will prove below that in this case the current  $\mathcal{J}$  is indeed equal to  $\mathcal{J}_i$ . If the flow value is smaller than  $F_G$ , then there must exist an edge  $(v_j, v_0)$  into the sink carrying less than  $m_j |I_j|$  units of flow. Hence there must exist an edge  $(u_k, v_j)$  carrying less than  $|I_j|$  units of flow because at least  $m_j$  jobs of  $\mathcal{J}$  are active in  $I_j$ . The algorithm removes the corresponding job  $J_k$ , i.e.  $\mathcal{J} := \mathcal{J} \setminus \{J_k\}$ . We will also prove below that this removal is correct, i.e.  $\mathcal{J}_i$  does not contain  $J_k$ . For the new set  $\mathcal{J}$  the algorithm starts the next round, invokes a maximum flow computation in the updated graph  $G(\mathcal{J}, \bar{m}, s)$ , where  $\bar{m}$  and  $s$  are updated too. The sequence of rounds ends when finally a flow of value  $F_G$  in the current graph  $G(\mathcal{J}, \bar{m}, s)$  is found. A formal description of the entire algorithm is given in Fig. 2. The pseudo-code uses a boolean variable set\_found that keeps track whether or not the desired  $\mathcal{J}_i$  has been correctly determined.

**Analysis of the algorithm** It remains to prove correctness of our algorithm. We prove that in any phase  $i$ ,  $\mathcal{J}_i$  is correctly determined. Lemma 3 immediately implies that the number  $m_{ij}$  of reserved processors, for  $1 \leq j < |\mathcal{I}|$ , as specified in lines 6 and 11 of the algorithm is correct. A schedule for  $\mathcal{J}_i$  on the reserved processors can be derived easily from the maximum flow computed in line 8. We described the construction of this schedule above when arguing that a flow of maximum value  $F_G$  in  $G(\mathcal{J}, \bar{m}, s)$  implies a feasible schedule using speed  $s$  on the reserved processors  $\bar{m}$ . It remains to show that sets

$\mathcal{J}_1, \dots, \mathcal{J}_p$  are correct. The proof is by induction on  $i$ . Consider an  $i$ ,  $1 \leq i \leq p$ , and suppose that  $\mathcal{J}_1, \dots, \mathcal{J}_{i-1}$  have been computed correctly. The inductive step and hence correctness of  $\mathcal{J}_i$  follow from [Lemmas 4 and 5](#).

**Lemma 4.** *In phase  $i$ , the set  $\mathcal{J}$  maintained by the algorithm always satisfies the invariant  $\mathcal{J}_i \subseteq \mathcal{J}$ .*

**Proof.** At the beginning of the phase  $\mathcal{J} = \{J_1, \dots, J_n\} \setminus (\mathcal{J}_1, \dots, \mathcal{J}_{i-1})$  is the set of all remaining jobs and the invariant is satisfied. Consider a round, represented by lines 5 to 10 of the algorithm, where in the beginning  $\mathcal{J}_i \subseteq \mathcal{J}$  holds and a job  $J_{k_0}$  is removed in line 10 at the end of the round. We prove that  $J_{k_0}$  does not belong to  $\mathcal{J}_i$ .

Consider the flow computed in line 8 of the round. We call a vertex  $v_j$  *saturated* if the amount of flow routed along the outgoing edge  $(v_j, v_0)$  is equal to the capacity  $m_j|I_j|$ . Otherwise  $v_j$  is *unsaturated*. Since the computed flow has value  $F < F_G = W/s = P = \sum_{1 \leq j < |I|} m_j|I_j|$ , there exists at least one unsaturated vertex. We now modify the flow so as to increase the number of unsaturated vertices. For an edge  $e$  in the graph, let  $f(e)$  be the amount of flow routed along it. The modification is as follows: While there exists a vertex  $u_k$  with flow  $f(u_k, v_j) < |I_j|$  into an unsaturated vertex  $v_j$  and positive flow  $f(u_k, v_{j'}) > 0$  into a saturated vertex  $v_{j'}$ , change the flow as follows. Along  $(u_k, v_j)$  and  $(v_j, v_0)$  we increase the flow by  $\epsilon$  and along  $(u_k, v_{j'})$  and  $(v_{j'}, v_0)$  we reduce the flow by  $\epsilon$ , where  $\epsilon = \frac{1}{2} \min\{f(u_k, v_{j'}), |I_j| - f(u_k, v_j), m_j|I_j| - f(v_j, v_0)\}$ . At  $u_k$ ,  $v_j$  and  $v_{j'}$  the flow conservation law is observed. By the choice of  $\epsilon$ , the new flow along  $(u_k, v_{j'})$  and  $(v_{j'}, v_0)$  is positive. Also, by the choice of  $\epsilon$ , the capacity constraints on  $(u_k, v_j)$  and  $(v_j, v_0)$  are not violated. Hence the new flow is feasible. Moreover, the total value  $F$  of the flow does not change. Vertex  $v_{j'}$  is a new unsaturated vertex while  $v_j$  continues to be unsaturated because  $\epsilon < m_j|I_j| - f(v_j, v_0)$  and hence the new flow along  $(v_j, v_0)$  is strictly smaller than the capacity. Also, since  $\epsilon < |I_j| - f(u_k, v_j)$ , the new flow along  $(u_k, v_j)$  remains strictly below  $|I_j|$ . The modifications stop when the required conditions are not satisfied anymore, which happens after less than  $|I|$  steps.

Given the modified flow, let  $U = \{j \mid v_j \text{ is unsaturated}\}$  be the indices of the unsaturated vertices. Intuitively, these vertices correspond to the intervals in which less than  $m_j|I_j|$  units of processing time would be scheduled. Let  $\mathcal{J}'$ ,  $\mathcal{J}'' \subseteq \mathcal{J}$ , be the set of jobs  $J_k$  such that  $f(u_0, u_k) = w_k/s$  and the outgoing edges leaving  $u_k$  that carry positive flow all lead into unsaturated vertices. We argue that the job  $J_{k_0}$  removed from  $\mathcal{J}$  in line 10 of the algorithm belongs to  $\mathcal{J}'$ : Prior to the flow modifications, there was an unsaturated vertex  $v_{j_0}$  having an incoming edge  $(u_{k_0}, v_{j_0})$  with  $f(u_{k_0}, v_{j_0}) < |I_{j_0}|$ . Throughout the flow modifications,  $v_{j_0}$  remains an unsaturated vertex, and any flow update on  $(u_{k_0}, v_{j_0})$  ensures that the amount of flow is strictly below  $|I_{j_0}|$ . Hence, if  $u_{k_0}$  had an outgoing edge with positive flow into a saturated vertex, the flow modifications would not have stopped. Moreover, if  $f(u_0, u_{k_0}) < w_{k_0}/s$  held, we could increase the flow along edges  $(u_0, u_{k_0})$ ,  $(u_{k_0}, v_{j_0})$  and  $(v_{j_0}, v_0)$ , thereby increasing the total value  $F$  of the flow.

Next let  $\mathcal{J}'$ ,  $\mathcal{J}'' \subseteq \mathcal{J} \setminus \mathcal{J}'$ , be the set of jobs  $J_k$  such that  $u_k$  has at least one edge  $(u_k, v_j)$  into an unsaturated vertex  $v_j$ . Here we argue that for any such job, all the edges  $(u_k, v_j)$  into unsaturated vertices  $v_j$  carry a flow of exactly  $|I_j|$  units: First observe that a job  $J_k \in \mathcal{J}''$  does not belong to  $\mathcal{J}'$  and hence (a) there is an edge  $(u_k, v_j)$  with positive flow into a saturated vertex  $v_j$  or (b)  $f(u_0, u_k) < w_k/s$ . In case (a), if there was an edge  $(u_k, v_j)$  carrying less than  $|I_j|$  units of flow into an unsaturated vertex  $v_j$ , the flow modifications would not have stopped. In case (b), we could increase the flow along  $(u_0, u_k)$ ,  $(u_k, v_j)$  and  $(v_j, v_0)$ , thereby raising the flow value  $F$ .

For any  $I_j$ , let  $\mathcal{J}'(I_j)$ , be the set of jobs of  $\mathcal{J}'$  that are active in  $I_j$ . Set  $\mathcal{J}''(I_j)$  is defined analogously. Given the modified flow, let  $t_{kj}$  be the amount of flow along  $(u_k, v_j)$ . Note that for no job  $J_k \in \mathcal{J} \setminus (\mathcal{J}' \cup \mathcal{J}'')$  there exists an edge  $(u_k, v_j)$  into an unsaturated vertex and hence  $J_k$  is not active in an  $I_j$  with  $j \in U$ . Hence, for any  $j \in U$ , since  $v_j$  is unsaturated,  $\sum_{J_k \in \mathcal{J}'(I_j) \cup \mathcal{J}''(I_j)} t_{kj} < m_j|I_j|$ . As  $t_{kj} = |I_j|$ , for  $J_k \in \mathcal{J}''(I_j)$ ,

$$\sum_{J_k \in \mathcal{J}'(I_j)} t_{kj} < (m_j - |\mathcal{J}''(I_j)|)|I_j|. \quad (1)$$

We next prove that the job  $J_{k_0}$  removed at the end of the round in line 10 of the algorithm does not belong to  $\mathcal{J}' \cap \mathcal{J}_i$  and hence does not belong to  $\mathcal{J}_i$  because, as argued above,  $J_{k_0} \in \mathcal{J}'$ . So suppose that  $J_{k_0} \in \mathcal{J}' \cap \mathcal{J}_i$ . We show that in this case the total processing volume of  $\mathcal{J}_i$  is not large enough to fill the reserved processing intervals using a continuous speed of  $s_i$ . We observe that  $s_i$  is at least as large as the speed  $s$  computed in the current round of the algorithm: If  $s_i < s$ , then in the optimal schedule, all jobs of  $\mathcal{J}$  would be processed at a speed smaller than  $s$ . However, the total length of the available processing intervals for  $\mathcal{J}$  is not more than  $P = \sum_{1 \leq j < |I|} m_j|I_j|$  and total work to be completed is  $W = \sum_{J_k \in \mathcal{J}} w_k$ . Hence a speed of  $s_i < s = W/P$  is not sufficient to finish the jobs in time.

Let  $m_{ij}$  be the number of processors used by  $\mathcal{J}_i$  in interval  $I_j$  in an optimal solution. We prove that the jobs of  $\mathcal{J}_i$  cannot fill all the reserved processors in intervals  $I_j$  with  $j \in U$  using a speed of  $s_i \geq s$ . Recall that no job of  $\mathcal{J} \setminus (\mathcal{J}' \cup \mathcal{J}'')$  is active in any interval  $I_j$  with  $j \in U$ . Only the jobs of  $\mathcal{J}'(I_j) \cap \mathcal{J}_i$  and  $\mathcal{J}''(I_j) \cap \mathcal{J}_i$  are active. The jobs of  $\mathcal{J}'(I_j) \cap \mathcal{J}_i$  must fill at least  $m_{ij} - |\mathcal{J}''(I_j) \cap \mathcal{J}_i|$  processors over a time window of length  $|I_j|$  because each job of  $\mathcal{J}''(I_j) \cap \mathcal{J}_i$  can fill at best one processor. Thus, over all intervals  $I_j$  with  $j \in U$ , the jobs of  $\mathcal{J}' \cap \mathcal{J}_i$  have to provide a work of at least

$$\sum_{j \in U} (m_{ij} - |\mathcal{J}''(I_j) \cap \mathcal{J}_i|)|I_j|s_i. \quad (2)$$

We show that this is not the case. The total work of  $\mathcal{J}' \cap \mathcal{J}_i$  is

$$\sum_{J_k \in \mathcal{J}' \cap \mathcal{J}_i} w_k = \sum_{J_k \in \mathcal{J}' \cap \mathcal{J}_i} f(u_0, u_k) s = \sum_{j \in U} \sum_{J_k \in \mathcal{J}' \cap \mathcal{J}_i} t_{kj} s, \quad (3)$$

where  $t_{kj}$  is again the amount of flow along  $(u_k, v_j)$  in the modified flow. The first and second equations hold because, for any  $J_k \in \mathcal{J}'$ ,  $f(u_0, u_k) = w_k/s$  and along edges  $(u_k, v_j)$  positive amounts of flow are routed only into unsaturated vertices  $v_j$ ,  $j \in U$ . For any  $j \in U$ , we have

$$\sum_{J_k \in \mathcal{J}'(I_j) \cap \mathcal{J}_i} t_{kj} \leq |\mathcal{J}'(I_j) \cap \mathcal{J}_i| |I_j|$$

because the edge capacity of  $(u_k, v_j)$  is  $|I_j|$ . For the unsaturated vertex  $v_{j_0}$  and the job  $J_{k_0}$ , determined in line 10 of the algorithm, the flow along  $(u_{k_0}, v_{j_0})$  is strictly below  $|I_{j_0}|$ . Hence for  $j_0 \in U$  the stronger relation

$$\sum_{J_k \in \mathcal{J}'(I_{j_0}) \cap \mathcal{J}_i} t_{kj_0} < |\mathcal{J}'(I_{j_0}) \cap \mathcal{J}_i| |I_{j_0}|$$

holds because by assumption  $J_{k_0} \in \mathcal{J}' \cap \mathcal{J}_i$ . Taking into account (1), we obtain that expression (3) is

$$\sum_{J_k \in \mathcal{J}' \cap \mathcal{J}_i} w_k < \sum_{j \in U} \min\{|\mathcal{J}'(I_j) \cap \mathcal{J}_i|, m_j - |\mathcal{J}''(I_j)|\} |I_j| s.$$

We will show that for any  $j \in U$ ,

$$\min\{|\mathcal{J}'(I_j) \cap \mathcal{J}_i|, m_j - |\mathcal{J}''(I_j)|\} \leq m_{ij} - |\mathcal{J}''(I_j) \cap \mathcal{J}_i|. \quad (4)$$

Since  $s_i \geq s$  this implies that the total work of  $\mathcal{J}' \cap \mathcal{J}_i$  is strictly smaller than  $\sum_{j \in U} (m_{ij} - |\mathcal{J}''(I_j) \cap \mathcal{J}_i|) |I_j| s_i$ , the expression in (2).

So suppose that (4) is violated for some  $j \in U$ . Then  $|\mathcal{J}'(I_j) \cap \mathcal{J}_i| > m_{ij} - |\mathcal{J}''(I_j) \cap \mathcal{J}_i|$  and the number of jobs of  $\mathcal{J}_i$  that are active in  $I_j$  is strictly larger than  $m_{ij}$ . Hence  $m_{ij}$  is the total number of processors not yet occupied by  $\mathcal{J}_1 \cup \dots \cup \mathcal{J}_{i-1}$ . As the number of jobs of  $\mathcal{J}$  active in  $I_j$  is at least as large as the corresponding number of jobs of  $\mathcal{J}_i$ , we have  $m_j = m_{ij}$ . Since  $|\mathcal{J}''(I_j)| \geq |\mathcal{J}''(I_j) \cap \mathcal{J}_i|$  we have  $m_j - |\mathcal{J}''(I_j)| \leq m_{ij} - |\mathcal{J}''(I_j) \cap \mathcal{J}_i|$ , which contradicts the assumption that (4) was violated.  $\square$

**Lemma 5.** When phase  $i$  ends,  $\mathcal{J} = \mathcal{J}_i$ .

**Proof.** When phase  $i$  ends, consider set  $\mathcal{J}$  and the flow in the current graph  $G = (\mathcal{J}, \vec{m}, s)$ . Again, for any edge  $e$  let  $f(e)$  be the amount of flow along  $e$ . Since the value of the computed flow is equal to  $F_G$  we have  $f(u_0, u_k) = w_k/s$ , for any  $J_k \in \mathcal{J}$ . By Lemma 4,  $\mathcal{J}_i \subseteq \mathcal{J}$ . So suppose  $\mathcal{J}_i \neq \mathcal{J}$ , which implies  $\mathcal{J}_i \subset \mathcal{J}$ . We prove that in this case the speed  $s_i$  used for  $\mathcal{J}_i$  in an optimal schedule satisfies  $s_i \leq s$  and then derive a contradiction.

In the graph  $G = (\mathcal{J}, \vec{m}, s)$  we remove the flow associated with any  $J_k \in \mathcal{J} \setminus \mathcal{J}_i$ . More specifically, for any such  $J_k$ , we remove  $f(u_k, v_j)$  units of flow along the path  $(u_0, u_k)$ ,  $(u_k, v_j)$  and  $(v_j, v_0)$  until the flow along  $(u_0, u_k)$  is zero. Let  $f'(v_j, v_0)$  be the new amount of flow along  $(v_j, v_0)$  after the modifications. Then for any interval  $I_j$ , at least  $\lceil f'(v_j, v_0)/|I_j| \rceil$  jobs of  $\mathcal{J}_i$  are active in  $I_j$  because for each job at most  $|I_j|$  units of flow are routed into  $v_j$ . Also, at least  $\lceil f'(v_j, v_0)/|I_j| \rceil$  processors are available because  $m_j = f(v_j, v_0)/|I_j|$  processors are available in  $I_j$ . Hence

$$s_i \leq \sum_{J_k \in \mathcal{J}_i} w_k / \sum_{1 \leq j < |I|} \lceil f'(v_j, v_0)/|I_j| \rceil |I_j|.$$

The latter expression is upper bounded by  $s$  because

$$\sum_{J_k \in \mathcal{J}_i} w_k = s \sum_{J_k \in \mathcal{J}_i} w_k/s = s \sum_{J_k \in \mathcal{J}_i} f(u_0, u_k) = s \sum_{1 \leq j < |I|} f'(v_j, v_0) \leq s \sum_{1 \leq j < |I|} \lceil f'(v_j, v_0)/|I_j| \rceil |I_j|.$$

Hence  $\mathcal{J}_i$  is processed at speed  $s_i \leq s$  while jobs of  $\mathcal{J} \setminus \mathcal{J}_i$  are processed at speed  $s'$  with  $s' < s_i \leq s$ . However, this is impossible because the minimum average speed to process the jobs of  $\mathcal{J}$  is  $s = \sum_{J_k \in \mathcal{J}} w_k / \sum_{1 \leq j < |I|} m_j |I_j|$ .  $\square$

We conclude with the following theorem.

**Theorem 1.** An optimal schedule can be computed in polynomial time using combinatorial techniques.



### 3. Online algorithms

We present adaptations and extensions of the single processor algorithms *Optimal Available* and *Average Rate* [15]. Throughout this section, we consider power functions of the form  $P(s) = s^\alpha$ , where  $\alpha > 1$ .

#### 3.1. Algorithm *Optimal Available*

For single processor environments, *Optimal Available* (OA) works as follows: Whenever a new job arrives, OA computes an optimal schedule for the currently available, unfinished jobs. While it is straightforward to extend this strategy to parallel processing environments, the corresponding competitive analysis becomes considerably more involved. We have to prove a series of properties of the online algorithm's schedule and analyze how a schedule changes in response to the arrival of a new job. Our algorithm for parallel processors works as follows.

**Algorithm OA(m).** Whenever a new job arrives, compute an optimal schedule for the currently available unfinished jobs. This can be done using the algorithm of Section 2.

At any given time  $t_0$ , let  $S_{OA(m)}$  denote the schedule of OA(m) for the currently available unfinished jobs. Let  $J_1, \dots, J_n$  be the corresponding jobs, which have arrived by time  $t_0$  but are still unfinished. Let  $w_i$  be the remaining processing volume of  $J_i$ ,  $1 \leq i \leq n$ . Since  $J_1, \dots, J_n$  are available at time  $t_0$  we can ignore release times and only have to consider the deadlines  $d_1, \dots, d_n$ . Let  $\mathcal{I} = \{t_0\} \cup \{d_i \mid 1 \leq i \leq n\}$  be the set of relevant time points and  $\tau_1 < \dots < \tau_{|\mathcal{I}|}$  be the sorted order of the elements of  $\mathcal{I}$ , where  $|\mathcal{I}| \leq n + 1$ . In the time horizon  $[\tau_1, \tau_{|\mathcal{I}|}]$ , the  $j$ -th interval is  $I_j = [\tau_j, \tau_{j+1}]$ , where  $1 \leq j < |\mathcal{I}|$ . Schedule  $S_{OA(m)}$  is an optimal schedule for the current jobs  $J_1, \dots, J_n$ ; otherwise when last computing a schedule, OA(m) would have obtained a better solution for time horizon  $[\tau_1, \tau_{|\mathcal{I}|}]$ . By Lemma 1 we always assume that in  $S_{OA(m)}$  each job is processed at a constant speed.

For the given time  $t_0$ , let  $S_{OPT}$  denote the set of optimal schedules for the unfinished jobs  $J_1, \dots, J_n$ . We have  $S_{OA(m)} \in S_{OPT}$ . Since  $P(s) = s^\alpha$ , there holds  $S_{OPT} = S'_{OPT}$ . Recall that the latter set was introduced in the beginning of Section 2 as the subset of optimal schedules that minimize energy consumption according to the power function  $P(s) = s^\alpha$ . Hence  $S_{OA(m)} \in S'_{OPT}$  and, by Lemma 2, in each interval  $I_j$  of  $S_{OA(m)}$  we can rearrange the schedule such that each processor uses a fixed, non-varying speed in  $I_j$ ,  $1 \leq j < |\mathcal{I}|$ . In the following we restrict ourselves to schedules  $S_{OA(m)}$  satisfying this property. Such schedules are feasible and incur a minimum energy. Let  $s_{l,j}$  be the speed used by processor  $l$  in  $I_j$ , where  $1 \leq l \leq m$  and  $1 \leq j < |\mathcal{I}|$ . Also note that Lemma 3 holds.

The next Lemma 6 states that we can modify  $S_{OA(m)}$  even further so that on each processor the speed levels used form a non-increasing sequence. The modification is obtained by simply permuting within each interval  $I_j$  the schedules on the  $m$  processors,  $1 \leq j < |\mathcal{I}|$ . Hence feasibility and energy consumption of the overall schedule are not affected. Throughout the analysis we always consider schedules  $S_{OA(m)}$  that also fulfill this additional property. This will be essential in the proofs of the subsequent Lemmas 7–9.

**Lemma 6.** Given  $S_{OA(m)}$ , in each interval  $I_j$ ,  $1 \leq j < |\mathcal{I}|$ , we can permute the schedules on the  $m$  processors such that the following property holds. For any processor  $l$ ,  $1 \leq l \leq m$ , inequality  $s_{l,j} \geq s_{l,j+1}$  is satisfied for  $j = 1, \dots, |\mathcal{I}| - 2$ .

**Proof.** In each interval  $I_j$ ,  $1 \leq j < |\mathcal{I}|$ , we simply sort the schedules on the  $m$  processors in order of non-increasing speeds such that  $s_{1,j} \geq \dots \geq s_{m,j}$ . Consider an arbitrary  $j$ ,  $1 \leq j \leq |\mathcal{I}| - 2$ , and suppose that the desired property does not hold for all  $l$ ,  $1 \leq l \leq m$ . Then let  $l_0$  be the smallest index such that  $s_{l_0,j} < s_{l_0,j+1}$ . In  $I_{j+1}$  the first  $l_0$  processors, running at speeds  $s_{1,j+1} \geq \dots \geq s_{l_0,j+1}$ , execute a set  $\mathcal{J}'(I_{j+1})$  of at least  $l_0$  different jobs in the interval because a job cannot be executed in parallel on two or more processors. The jobs of  $\mathcal{J}'(I_{j+1})$  can only be processed on the first  $l_0 - 1$  processors in  $I_j$  because in  $S_{OA(m)}$  each job is processed at a constant speed and the last  $m - l_0 + 1$  processors in  $I_j$  use speeds which are strictly smaller than  $s_{1,j+1}, \dots, s_{l_0,j+1}$ . Consider a time window  $W = [\tau_j, \tau_j + \delta)$  with  $\delta > 0$  in  $I_j$  such that each of the first  $l_0$  processors does not change the job it executes in  $W$ . Among the at least  $l_0$  jobs of  $\mathcal{J}'(I_{j+1})$  there must exist a  $J_i$  which is not processed in  $W$  because jobs of  $\mathcal{J}'(I_{j+1})$  are not executed on the  $m - l_0 + 1$  last processors in  $I_j$  and the first  $l_0 - 1$  processors handle at most  $l_0 - 1$  different jobs of  $\mathcal{J}'(I_{j+1})$  in  $W$ . Suppose that  $J_i$  is processed for  $\delta_i$  time units in  $I_{j+1}$  using speed  $s(J_i)$ . Set  $\delta' = \min\{\delta, \delta_i\}$  and  $\epsilon = (s_{l_0,j+1} - s_{l_0,j})/2$ . We now modify the schedule as follows. In  $I_{j+1}$  for  $\delta'$  time units where  $J_i$  is processed we reduce the speed by  $\epsilon$ . In  $I_j$ , in the time window  $W' = [\tau_j, \tau_j + \delta')$  we increase the speed of processor  $l_0$  by  $\epsilon$  and use the extra processing capacity of  $\delta'\epsilon$  units to process  $\delta'\epsilon$  units of  $J_i$ . We obtain a feasible schedule. Since  $s(J_i) - \epsilon \geq s_{j+1,l_0} - \epsilon = s_{j,l_0} + \epsilon$ , by the convexity of the power function  $P(s) = s^\alpha$  the modified schedule has a strictly smaller energy consumption. We obtain a contradiction to the fact that the original schedule was optimal.  $\square$

We investigate the event that a new job  $J_{n+1}$  arrives at time  $t_0$ . As usual  $d_{n+1}$  and  $w_{n+1}$  are the deadline and the processing volume of  $J_{n+1}$ . Let  $S_{OA(m)}$  be the schedule of OA(m) immediately before the arrival of  $J_{n+1}$ , and let  $S'_{OA(m)}$  be the schedule immediately after the arrival when OA(m) has just computed a new schedule. We present two lemmas that relate the two schedules. Loosely speaking we prove that processor speeds can only increase. The speed at which a

job  $J_k$ ,  $1 \leq k \leq n$ , is processed in  $S'_{OA(m)}$  is at least as high as the corresponding speed in  $S_{OA(m)}$ . Furthermore, at any time  $t \geq t_0$ , the minimum processor speed in  $S'_{OA(m)}$  is at least as high as that in  $S_{OA(m)}$ .

Schedules  $S_{OA(m)}$  and  $S'_{OA(m)}$  are optimal schedules for the respective workloads. Recall that they process jobs at constant speeds. Let  $s(J_k)$  be the speed at which  $J_k$  is processed in  $S_{OA(m)}$ ,  $1 \leq k \leq n$ . Let  $s'(J_k)$  be the speed used for  $J_k$  in  $S'_{OA(m)}$ ,  $1 \leq k \leq n+1$ .

**Lemma 7.** *There holds  $s'(J_k) \geq s(J_k)$ , for any  $1 \leq k \leq n$ .*

**Proof.** In  $S_{OA(m)}$  let  $s_1 > s_2 > \dots > s_p$  be the speeds used in the schedule and let  $\mathcal{J}_1, \dots, \mathcal{J}_p$  be the associated job sets, i.e. the jobs of  $\mathcal{J}_i$  are processed at speed  $s_i$ ,  $1 \leq i \leq p$ . In the following we assume that the statement of the lemma does not hold and derive a contradiction. So let  $i$  be the smallest index such that  $\mathcal{J}_i$  contains a job  $J_k$  with  $s'(J_k) < s(J_k) = s_i$ . We partition  $\mathcal{J}_i$  into two sets  $\mathcal{J}_{i1}$  and  $\mathcal{J}_{i2}$  such that all  $J_k \in \mathcal{J}_{i1}$  satisfy  $s'(J_k) \geq s_i$ , as desired. For each job  $J_k \in \mathcal{J}_{i2}$  we have  $s'(J_k) < s_i$ . By the choice of  $i$ , any job  $J_k \in \mathcal{J}_1 \cup \dots \cup \mathcal{J}_{i-1}$  also fulfills the desired property  $s'(J_k) \geq s(J_k)$ .

Let  $t_1$  be the first point of time when the processing of any job of  $\mathcal{J}_{i2}$  starts in  $S'_{OA(m)}$ . Let  $P$  be a processor executing such a job at time  $t_1$ . We first argue that any  $J_k \in \mathcal{J}_1 \cup \dots \cup \mathcal{J}_{i-1} \cup \mathcal{J}_{i1}$  that is active at time  $t \geq t_1$ , i.e. that satisfies  $d_k > t_1$ , is scheduled at all times throughout  $[t_1, d_k]$ : So suppose  $J_k$  were not executed on any of the processors in a time window  $W \subseteq [t_1, d_k]$  and let  $\delta > 0$  be the length of  $W$ . Determine a  $\delta'$  with  $0 < \delta' \leq \delta$  such that during the first  $\delta'$  time units of  $W$  processor  $P$  does not change the job it executes. We set  $\delta_k = \min\{\delta', w_k/s'(J_k)\}$ . Here  $w_k/s'(J_k)$  is the total time for which  $J_k$  is scheduled in  $S'_{OA(m)}$ . Let  $s(P, W)$  be the speed used by processor  $P$  within the first  $\delta'$  time units of  $W$  and set  $\epsilon = (s'(J_k) - s(P, W))/2$ . The value of  $\epsilon$  is strictly positive because  $s'(J_k) \geq s_i > s(P, W)$ . The last inequality holds because during the first  $\delta'$  time units of  $W$  processor  $P$  executes a job of  $\mathcal{J}_{i2}$  at a speed that is strictly smaller than  $s_i$ . We now modify  $S'_{OA(m)}$  as follows. During any  $\delta_k$  time units where  $J_k$  is processed we reduce the speed by  $\epsilon$ . During the first  $\delta_k$  time units of  $W$  we increase the speed by  $\epsilon$  and use the extra processing capacity of  $\delta_k \epsilon$  units to execute  $\delta_k \epsilon$  units of  $J_k$ . Hence we obtain a new feasible schedule. Since  $s'(J_k) - \epsilon = s(P, W) + \epsilon$ , by the convexity of the power function  $P(s) = s^\alpha$ , the schedule has a strictly smaller energy consumption, contradicting the fact that  $S'_{OA(m)}$  is an optimal schedule for the given workload.

Next consider any fixed time  $t \geq t_1$ . In  $S_{OA(m)}$  let  $m_{i2}$  be the number of processors that execute jobs of  $\mathcal{J}_{i2}$  at time  $t$ . Similarly, in  $S'_{OA(m)}$  let  $m'_{i2}$  be the number of processors that execute jobs of  $\mathcal{J}_{i2}$  at time  $t$ . We will prove  $m_{i2} \geq m'_{i2}$ . Let  $n_l$ , where  $1 \leq l \leq i-1$ , be the number of jobs  $J_k$  of  $\mathcal{J}_l$  that are active at time  $t$ , i.e. that satisfy  $d_k > t$ . Furthermore, let  $n_{i1}$  and  $n_{i2}$  be the number of jobs of  $\mathcal{J}_{i1}$  and  $\mathcal{J}_{i2}$ , respectively, that are active at time  $t$ . As shown in the last paragraph, all jobs of  $\mathcal{J}_1 \cup \dots \cup \mathcal{J}_{i-1} \cup \mathcal{J}_{i1}$  that are active at time  $t$  are also processed at this time in  $S'_{OA(m)}$ . Hence

$$m'_{i2} \leq m - \sum_{l=1}^{i-1} n_l - n_{i1}. \quad (5)$$

Strict inequality holds, for instance, if the new job  $J_{n+1}$  or a job of  $\mathcal{J}_{i+1} \cup \dots \cup \mathcal{J}_p$  is scheduled at time  $t$ . Let  $m_l$  be the number of processors executing jobs of  $\mathcal{J}_l$  in  $S_{OA(m)}$ , where  $1 \leq l \leq i$ . By Lemma 3,  $m_i = \min\{n_i, m - \sum_{l=1}^{i-1} m_l\}$ . We remark that Lemma 3 is formulated with respect to an interval  $I_j$ . We can apply the statement simply by considering the interval containing  $t$ . Since  $m_l \leq n_l$ , for  $1 \leq l \leq i-1$ , we obtain

$$m_i \geq \min\{n_i, m - \sum_{l=1}^{i-1} n_l\}. \quad (6)$$

If the latter minimum is given by the term  $m - \sum_{l=1}^{i-1} n_l$ , then in (5) we obtain  $m'_{i2} \leq m - \sum_{l=1}^{i-1} n_l - n_{i1} \leq m_i - n_{i1} \leq m_{i2}$ . The last inequality holds because  $m_{i2}$  cannot be smaller than  $m_i$ , the number of machines used by jobs of  $\mathcal{J}_i$ , minus the number of jobs of  $\mathcal{J}_{i1}$  active at time  $t$ . Hence  $m'_{i2} \leq m_{i2}$ , as desired. If the minimum in (6) is given by  $n_i$ , then since  $m_{i2} \geq m_i - n_{i1}$ , we have  $m_{i2} \geq n_i - n_{i1}$ . As  $n_i - n_{i1} = n_{i2} \geq m'_{i2}$  we conclude again  $m_{i2} \geq m'_{i2}$ .

Finally let  $T'$  be the total processing time used for jobs of  $\mathcal{J}_{i2}$  in  $S'_{OA(m)}$ . This value can be determined as follows. The time horizon of  $S'_{OA(m)}$  is partitioned into maximal intervals such that the number of processors handling jobs of  $\mathcal{J}_{i2}$  does not change within an interval. For each such interval, the length of the interval is multiplied by the number of processors executing jobs of  $\mathcal{J}_{i2}$ . Value  $T'$  is simply the sum of all these products. Similarly, let  $T$  be the total processing time used for jobs of  $\mathcal{J}_{i2}$  in  $S_{OA(m)}$ . At any time  $t$  we have  $m_{i2} \geq m'_{i2}$ . This was shown above for times  $t \geq t_1$ . It trivially holds for times  $t$  with  $t_0 \leq t < t_1$  because no jobs of  $\mathcal{J}_{i2}$  are processed in  $S'_{OA(m)}$  at that time. Hence the number of processors handling jobs of  $\mathcal{J}_{i2}$  in  $S_{OA(m)}$  is always at least as high as the corresponding number in  $S'_{OA(m)}$ . This implies  $T \geq T'$ . The total processing volume  $\sum_{J_k \in \mathcal{J}_{i2}} w_k$  of jobs in  $\mathcal{J}_{i2}$  fills  $T'$  time units using a speed or speeds of  $s < s_i$  in  $S'_{OA(m)}$ . Thus using a speed of  $s_i$  the processing volume is not sufficient to fill  $T > T'$  time units in  $S_{OA(m)}$ . We obtain a contradiction to the fact that jobs of  $\mathcal{J}_i$  are processed at constant speed  $s_i$  in  $S_{OA(m)}$ .  $\square$

**Lemma 8.** *At any time  $t$ , where  $t \geq t_0$ , the minimum speed used by the  $m$  processors at time  $t$  in  $S'_{OA(m)}$  is at least as high as the minimum speed used on the  $m$  processors at time  $t$  in  $S_{OA(m)}$ .*

**Proof.** Consider an arbitrary time  $t \geq t_0$ . Let  $s_{\min}$  be the minimum speed on the  $m$  processors in  $S_{OA(m)}$  at time  $t$ . Let  $s'_{\min}$  be the corresponding minimum speed in  $S'_{OA(m)}$  at time  $t$ . If  $s_{\min} = 0$ , then the lemma trivially holds. So suppose  $s_{\min} > 0$  and assume that the statement of the lemma does not hold, i.e.  $s'_{\min} < s_{\min}$ . Since  $s_{\min} > 0$ ,  $m$  different jobs are processed at time  $t$  in  $S_{OA(m)}$ . By Lemma 7, for any job, the speed at which it is processed in  $S'_{OA(m)}$  is at least as high as the corresponding speed in  $S_{OA(m)}$ . Hence among the  $m$  jobs processed in  $S_{OA(m)}$  at time  $t$  there must exist at least one job  $J_k$  that is not executed in  $S'_{OA(m)}$  at time  $t$ . For this job we have  $d_k > t$  because  $J_k$  is available for processing and hence active at time  $t$ . Determine a  $\delta_k > 0$  with  $\delta_k \leq \min\{d_k - t, w_k/s'(J_k)\}$  such that  $J_k$  is not processed in the interval  $[t, t + \delta_k)$  in  $S'_{OA(m)}$ . Choose a  $\delta$  with  $0 < \delta \leq \delta_k$  such that in  $S'_{OA(m)}$  the last processor  $m$ , running at minimum speed  $s'_{\min}$  at time  $t$ , does not change the job it executes within time window  $W = [t, t + \delta)$ . Throughout  $W$  a speed of  $s'_{\min}$  is used on processor  $m$ . Set  $\epsilon = (s_{\min} - s'_{\min})/2$ , which is a strictly positive value.

We now modify  $S'_{OA(m)}$ . During  $\delta$  time units in which  $J_k$  is processed we reduce the speed by  $\epsilon$ . These time units exist because  $\delta \leq w_k/s'(J_k)$ . In the time window  $W$  we increase the speed on the last processor  $m$  by  $\epsilon$  and use the extra processing capacity of  $\delta\epsilon$  unit to process  $\delta\epsilon$  units of  $J_k$ . Since  $\delta \leq d_k - t$  job  $J_k$  finishes by its deadline, and we obtain a feasible schedule. Throughout  $W$  the newly set speed on the last processor  $m$  is  $s'_{\min} + \epsilon$ . Finally notice  $s'_{\min} + \epsilon = s_{\min} - \epsilon \leq s(J_k) - \epsilon \leq s'(J_k) - \epsilon$ . The last inequality follows from Lemma 7. Hence  $s'_{\min} + \epsilon \leq s'(J_k) - \epsilon$ , and by the convexity of the power function  $P(s) = s^\alpha$ , we obtain a schedule with a strictly smaller power consumption. This is a contradiction to the fact that  $S'_{OA(m)}$  is optimal.  $\square$

We finally need a lemma that specifies the minimum processor speed for a time window if a job finishes at some time strictly before its deadline.

**Lemma 9.** *If in  $S_{OA(m)}$  a job  $J_k$  finishes at some time  $t$  with  $t < d_k$ , then throughout  $[t, d_k)$  the minimum speed on the  $m$  processors is at least  $s(J_k)$ .*

**Proof.** Consider a job  $J_k$  that finishes at time  $t < d_k$ . Suppose that the lemma does not hold and let  $t'$  with  $t \leq t' < d_k$  be a time such that the minimum speed on the  $m$  processors is strictly smaller than  $s(J_k)$ . Recall that we consider schedules satisfying the property of Lemma 6, i.e. on any processor the speed levels form a non-increasing sequence. Hence within  $W = [t', d_k)$  there exists a processor  $P$  whose speed throughout  $W$  is upper bounded by some  $s < s(J_k)$ . Choose a positive  $\delta$  with  $\delta \leq \min\{d_k - t', w_k/s(J_k)\}$  such that processor  $P$  does not switch jobs in  $W' = [t', t' + \delta)$ . Set  $\epsilon = (s(J_k) - s)/2$ . We modify the schedule as follows. During  $\delta$  time units where  $J_k$  is processed we reduce the speed by  $\epsilon$ . In  $W'$  we increase the speed of processor  $P$  by  $\epsilon$  and use the extra processing capacity to finish the missing  $\delta\epsilon$  processing units of  $J_k$ . The resulting schedule is feasible and incurs a strictly smaller energy consumption. Hence we obtain a contradiction to the fact that  $S_{OA(m)}$  is an optimal schedule.  $\square$

We next turn to the competitive analysis of  $OA(m)$ . We use a potential function that is inspired by a potential used by Bansal et al. [5]. We first define our modified potential function. At any time  $t$  consider the schedule of  $S_{OA(m)}$ . Let  $s_1 > \dots > s_p$  be the speeds used in the schedule and let  $\mathcal{J}_1, \dots, \mathcal{J}_p$  be the associated job sets, i.e. jobs of  $\mathcal{J}_i$  are processed at speed  $s_i$ . Let  $W_{OA(m)}(i)$  be the total remaining processing volume of jobs  $J_k \in \mathcal{J}_i$  in  $OA(m)$ 's schedule  $1 \leq i \leq p$ . Similarly, let  $W_{OPT}(i)$  be the total remaining processing volume that  $OPT$  has to finish for jobs  $J_k \in \mathcal{J}_i$ ,  $1 \leq i \leq p$ . We remark that, for a job  $J_k \in \mathcal{J}_i$ ,  $OPT$ 's remaining work might be larger than that of  $OA(m)$  if  $OPT$  has completed less work on  $J_k$  so far. Sets  $\mathcal{J}_1, \dots, \mathcal{J}_p$  represent the jobs that  $OA(m)$  has not yet finished. Let  $\mathcal{J}'$  be the set of jobs that are finished by  $OA(m)$  but still unfinished by  $OPT$ . Obviously, these jobs have a deadline after the current time  $t$ . If  $\mathcal{J}' \neq \emptyset$ , then for any job of  $\mathcal{J}'$  consider the speed  $OA(m)$  used when last processing the job. Partition  $\mathcal{J}'$  according to these speeds, i.e.  $\mathcal{J}'_i$  consists of those jobs that  $OA(m)$  executed last at speed  $s'_i$ , where  $1 \leq i \leq p'$  for some positive  $p'$ . Let  $W'_{OPT}(i)$  be the total remaining processing volume  $OPT$  has to finish on jobs of  $\mathcal{J}'_i$ ,  $1 \leq i \leq p'$ . If  $\mathcal{J}' = \emptyset$ , we simply set  $s'_1 = 0$  and  $\mathcal{J}'_1 = \emptyset$ . The potential  $\Phi$  is defined as follows.

$$\Phi = \alpha \sum_{i \geq 1} (s_i)^{\alpha-1} (W_{OA(m)}(i) - \alpha W_{OPT}(i)) - \alpha^2 \sum_{i \geq 1} (s'_i)^{\alpha-1} W'_{OPT}(i)$$

Compared to the potential function used to analyze  $OA$  in the single processor setting, our function here consists of a second term representing  $OPT$ 's unfinished work. This second term is essential to establish the competitiveness of  $\alpha^\alpha$  and, in particular, is crucially used in the analysis of the working case (see below) where we have to consider  $m$  processor pairs.

At any time  $t$  let  $s_{OA(m),1}(t) \geq \dots \geq s_{OA(m),m}(t)$  and  $s_{OPT,1}(t) \geq \dots \geq s_{OPT,m}(t)$  be the speeds used by  $OA(m)$  and  $OPT$  on the  $m$  processors, respectively. In the remainder of this section we will prove the following properties.

- (a) Whenever a new job arrives or a job is finished by  $OA(m)$  or  $OPT$ , the potential does not increase.  
 (b) At all other times

$$\sum_{l=1}^m (s_{OA(m),l}(t))^\alpha - \alpha^\alpha \sum_{l=1}^m (s_{OPT,l}(t))^\alpha + \frac{d\Phi(t)}{dt} \leq 0. \quad (7)$$

Integrating over all times we obtain that  $OA(m)$  is  $\alpha^\alpha$ -competitive.

**Working case** We prove property (b). So let  $t$  be any time when the algorithms are working on jobs. In a first step we match the processors of  $OA(m)$  to those of  $OPT$ . This matching is constructed as follows. First, processors handling the same jobs are paired. More specifically, if a job  $J_k$  is executed on a processor  $l$  in  $OA(m)$ 's schedule and on processor  $l'$  in  $OPT$ 's schedule, then we match the two processors  $l$  and  $l'$ . All remaining processors of  $OA(m)$  and  $OPT$ , handling disjoint job sets, are matched in an arbitrary way.

For each matched processor pair  $l, l'$ , we consider its contribution in (7) and prove that the contribution is non-positive. Summing over all pairs  $l, l'$ , this establishes (7). If the speed of processor  $l'$  in  $OPT$ 's schedule is 0, then the analysis is simple: If the speed on processor  $l$  in  $OA(m)$  schedule is also 0, then the contribution of the processor pair  $l, l'$  in (7) is 0. If the speed on processor  $l$  is positive, then  $OA(m)$  executes a job on the processor. Suppose that the job is contained in set  $\mathcal{J}_i$ . Then in  $\Phi$  the value of  $W_{OA(m)}(i)$  decreases at rate  $s_i$  and the contribution of the processor pair  $l, l'$  in (7) is  $s_i^\alpha - \alpha(s_i)^{\alpha-1}s_i < 0$ .

Hence in the following we assume that the speed of processor  $l'$  in  $OPT$ 's schedule is positive. We show that  $OA(m)$  also executes a job  $J_k$  on its processor  $l$  and that the contribution of the processor pair  $l, l'$  is upper bounded by

$$(1 - \alpha)s_i^\alpha + \alpha^2 s_i^{\alpha-1} s_{OPT,l'}(t) - \alpha^\alpha (s_{OPT,l'}(t))^\alpha, \quad (8)$$

where  $i$  is the index such that  $J_k \in \mathcal{J}_i$ . As in [5] one can then show the non-positivity of this expression. First consider the case that both processors  $l$  and  $l'$  handle the same job  $J_k$ . Let  $\mathcal{J}_i$  be the set containing  $J_k$ . Then  $W_{OA(m)}(i)$  decreases at rate  $s_i$  and  $W_{OPT}(i)$  decreases at rate  $s_{OPT,l'}(t)$ . Hence the contribution of the processor pair in (7) is  $s_i^\alpha - \alpha^\alpha (s_{OPT,l'}(t))^\alpha + (-\alpha(s_i)^{\alpha-1}s_i + \alpha^2(s_i)^{\alpha-1}s_{OPT,l'}(t)) = (1 - \alpha)s_i^\alpha + \alpha^2(s_i)^{\alpha-1}s_{OPT,l'}(t) - \alpha^\alpha (s_{OPT,l'}(t))^\alpha$ , as stated in (8).

Next consider the case that the job  $J_{k'}$  executed on processor  $l'$  in  $OPT$ 's schedule is not executed by  $OA(m)$  on processor  $l$  and hence is not executed on any of its processors at time  $t$ . If  $J_{k'}$  is still unfinished by  $OA(m)$ , then let  $\mathcal{J}_{i'}$  be the set containing  $J_{k'}$ . Recall that  $OA(m)$ 's schedule is optimal. Hence, by Lemma 3, it always processes the jobs from the smallest indexed sets  $\mathcal{J}_i$ , which in turn use the highest speeds  $s_i$ . Thus, if  $J_{k'}$  is not executed by  $OA(m)$ , then the schedule currently processes  $m$  jobs, each of which is contained in  $\mathcal{J}_1 \cup \dots \cup \mathcal{J}_{i'-1}$  or contained in  $\mathcal{J}_{i'}$  but is different from  $J_{k'}$ . Let  $J_k$  be the job executed by  $OA(m)$  on processor  $l$  and let  $\mathcal{J}_i$  be the job set containing  $J_k$ . Then  $i \leq i'$  and  $s_i \geq s_{i'}$ . In  $\Phi$  the term  $W_{OA(m)}(i)$  decreases at rate  $s_i$  while the term  $W_{OPT}(i')$  decreases at rate  $s_{OPT,l'}(t)$ . Hence the contribution of the processor pair  $l, l'$  in (7) is  $s_i^\alpha - \alpha^\alpha (s_{OPT,l'}(t))^\alpha + (-\alpha(s_i)^{\alpha-1}s_i + \alpha^2(s_{i'})^{\alpha-1}s_{OPT,l'}(t)) \leq (1 - \alpha)s_i^\alpha + \alpha^2(s_i)^{\alpha-1}s_{OPT,l'}(t) - \alpha^\alpha (s_{OPT,l'}(t))^\alpha$ , which is the bound of (8).

It remains to consider the case that  $J_{k'}$  is already finished by  $OA(m)$ . Then  $J_{k'}$  is contained in the set  $\mathcal{J}'$  of jobs that are finished by  $OA(m)$  but unfinished by  $OPT$ . Let  $\mathcal{J}_{i'}$  be the set containing  $J_{k'}$ . In  $\Phi$  the term  $W'_{OPT}(i')$  decreases at rate  $s_{OPT,l'}(t)$ . When  $OA(m)$  finished  $J_{k'}$  in its schedule, it used speed  $s'_{i'}$ . By Lemma 9, in the former schedule of  $OA(m)$ , the minimum processor speed throughout the time window until  $d_{k'}$  was always at least  $s'_{i'}$ . This time window contains the current time, i.e.  $t < d_{k'}$ . Since the completion of  $J_{k'}$   $OA(m)$  might have computed new schedules in response to the arrival of new jobs but, by Lemma 8, the minimum processor speed at any time can only increase. Hence in the current schedule of  $OA(m)$  and at the current time  $t$  the minimum processor speed is at least  $s'_{i'}$ . Thus on its processor  $l$   $OA(m)$  uses a speed of at least  $s'_{i'} > 0$ . Hence  $OA(m)$  executes a job  $J_k$  that belongs to, say, set  $\mathcal{J}_i$ . We have  $s_i \geq s'_{i'}$ . In  $\Phi$  the term  $W_{OA(m)}(i)$  decreases at rate  $s_i$ . We conclude again that the contribution of the processor pair  $l, l'$  is  $s_i^\alpha - \alpha^\alpha (s_{OPT,l'}(t))^\alpha + (-\alpha(s_i)^{\alpha-1}s_i + \alpha^2(s'_{i'})^{\alpha-1}s_{OPT,l'}(t)) \leq (1 - \alpha)s_i^\alpha + \alpha^2(s_i)^{\alpha-1}s_{OPT,l'}(t) - \alpha^\alpha (s_{OPT,l'}(t))^\alpha$ , which is again the bound of (8).

**Completion/arrival case** When  $OPT$  finishes a job  $J_k$  the potential does not change because in the terms  $W_{OPT}(i)$  or  $W_{OPT}(i')$  the remaining processing volume of  $J_k$  simply goes to 0. Similarly if  $OA(m)$  completes a job already finished by  $OPT$ , the potential does not change. If  $OA(m)$  completes a job  $J_k$  not yet finished by  $OPT$ , then let  $\mathcal{J}_i$  be the set containing  $J_k$  immediately before completion. In the first term of  $\Phi$ ,  $W_{OPT}(i)$  increases by  $\alpha^2(s_i)^{\alpha-1}w_k$ , where  $w_k$  is the remaining work of  $J_k$  to be done by  $OPT$ . However, at the same time, the second term of  $\Phi$  decreases by exactly  $\alpha^2(s_i)^{\alpha-1}w_k$  because  $J_k$  opens or joins a set  $\mathcal{J}'_{i'}$  with  $s'_{i'} = s_i$ .

When a new job  $J_{n+1}$  arrives, the structure of the analysis is similar to that in the single processor case, see [5]. We imagine that a job  $J_{n+1}$  of processing volume 0 arrives. Then we increase the processing volume to its true size. However, in the multi-processor setting, we have to study how the schedule of  $OA(m)$  changes if the processing volume of  $J_{n+1}$  increases from a value  $w_{n+1}$  to  $w'_{n+1} = w_{n+1} + \epsilon$ , for some  $\epsilon > 0$ . Let  $S_{OA(m)}$  and  $S'_{OA(m)}$  be the schedules of  $OA(m)$  before and after the increase, respectively. For any job  $J_k$ ,  $1 \leq k \leq n+1$ , let  $s(J_k)$  and  $s'(J_k)$  be the speeds used for  $J_k$  in  $S_{OA(m)}$  and  $S'_{OA(m)}$ , respectively. The next lemma states that the speeds at which jobs are processed does not decrease. The proof is identical to that of Lemma 7.

**Lemma 10.** *There holds  $s'(J_k) \geq s(J_k)$ , for any  $1 \leq k \leq n+1$ .*

In  $S_{OA(m)}$ , let  $\mathcal{J}_{i_0}$  be the job set containing  $J_{n+1}$ , i.e.  $J_{n+1} \in \mathcal{J}_{i_0}$ , where  $1 \leq i_0 \leq p$ . The following lemma implies that the speeds of jobs  $J_k \in \mathcal{J}_{i_0+1} \cup \dots \cup \mathcal{J}_p$  do not change when increasing the processing volume of  $J_{n+1}$ .

**Lemma 11.** *There holds  $s'(J_k) = s(J_k)$ , for any  $J_k \in \mathcal{J}_{i_0+1} \cup \dots \cup \mathcal{J}_p$ .*

**Proof.** We first analyze  $S_{OA(m)}$  and  $S'_{OA(m)}$  with respect to the times when jobs of  $\mathcal{J}_{i_0+1} \cup \dots \cup \mathcal{J}_p$  are processed. Consider any time  $t \geq t_0$  in the scheduling horizon, where  $t_0$  is again the current time. Let  $m_1$  be the number of processors executing jobs of  $\mathcal{J}_1 \cup \dots \cup \mathcal{J}_{i_0}$  in  $S_{OA(m)}$  at time  $t$ , and let  $m_2$  be the number of processors executing jobs of  $\mathcal{J}_{i_0+1} \cup \dots \cup \mathcal{J}_p$  in  $S_{OA(m)}$  at time  $t$ . Values  $m'_1$  and  $m'_2$  are defined analogously with respect to  $S'_{OA(m)}$ . In the following we will prove  $m'_2 \geq m_2$ .

If  $m_2 = 0$ , there is nothing to show. So suppose  $m_2 > 0$ . By Lemma 3 exactly  $m_1$  jobs of  $\mathcal{J}_1 \cup \dots \cup \mathcal{J}_{i_0}$  are active at time  $t$ , i.e. have a deadline after time  $t$ ; if there were more active jobs,  $OA(m)$  would have assigned more processors to the jobs of  $\mathcal{J}_1 \cup \dots \cup \mathcal{J}_{i_0}$  because they are executed at higher speeds than the jobs of  $\mathcal{J}_{i_0+1} \cup \dots \cup \mathcal{J}_p$ . We next assume  $m'_2 < m_2$  and derive a contradiction. We distinguish two cases depending on whether or not  $S'_{OA(m)}$  has an idle processor at time  $t$ .

If there is no idle processor in  $S'_{OA(m)}$  at time  $t$ , then  $m'_1 = m - m'_2 > m - m_2 \geq m_1$  jobs of  $\mathcal{J}_1 \cup \dots \cup \mathcal{J}_{i_0}$  must be executed in  $S'_{OA(m)}$  at time  $t$ . However, this is impossible because only  $m_1$  jobs of this union of sets are active and hence available for processing at time  $t$ . On the other hand, if there is a processor  $P$  that is idle in  $S'_{OA(m)}$  at time  $t$ , then there must exist a time window  $W$  containing time  $t$  such that  $P$  is idle throughout  $W$ . Moreover, since  $m'_2 < m_2$ , there must exist a job  $J_k \in \mathcal{J}_{i_0+1} \cup \dots \cup \mathcal{J}_p$  that is active but not processed at time  $t$ . Hence, within  $W$  we can select a time window  $W' \subseteq W$  of length  $\delta$ , where  $\delta \leq w(J_k)/s'(J_k)$ , such that  $J_k$  is not executed on any processor in  $S'_{OA(m)}$  throughout  $W'$ . We now modify  $S'_{OA(m)}$  as follows. For any  $\delta$  time units where  $J_k$  is scheduled, we reduce the speed to  $s'(J_k)/2$ . On processor  $P$  we execute  $J_k$  at speed  $s'(J_k)/2$  for  $\delta$  time units in  $W'$ . By convexity of the power function, we obtain a better schedule with a strictly smaller energy consumption, contradicting the fact that  $S'_{OA(m)}$  is an optimal schedule.

Let  $T$  be the total time used to process jobs of  $\mathcal{J}_{i_0+1} \cup \dots \cup \mathcal{J}_p$  in  $S_{OA(m)}$ . As in the proof of Lemma 7, this time is computed as follows. We partition the scheduling horizon of  $S_{OA(m)}$  into maximal intervals such that within each interval the number of processors executing jobs of  $\mathcal{J}_{i_0+1} \cup \dots \cup \mathcal{J}_p$  does not change. The length of each interval is multiplied by the number of processors handling jobs of  $\mathcal{J}_{i_0+1} \cup \dots \cup \mathcal{J}_p$ . Then  $T$  is simply the sum of these products. Analogously, let  $T'$  be the total time used to process jobs of  $\mathcal{J}_{i_0+1} \cup \dots \cup \mathcal{J}_p$  in  $S'_{OA(m)}$ . In the last paragraph we showed that at any time  $t$ ,  $m'_2 \geq m_2$ . Hence  $T' \geq T$ .

We are now ready to finish the proof of our lemma. By Lemma 10 we have  $s'(J_k) \geq s(J_k)$  for any  $J_k \in \mathcal{J}_{i_0+1} \cup \dots \cup \mathcal{J}_p$ . Suppose that in the latter set there exists a job  $J_{k_0}$  with  $s'(J_{k_0}) > s(J_{k_0})$ . In  $S_{OA(m)}$  the jobs of  $\mathcal{J}_{i_0+1} \cup \dots \cup \mathcal{J}_p$  are processed for  $T$  time units. The total processing volume of these jobs cannot fill  $T' \geq T$  time units if  $J_{k_0}$  is processed at a strictly higher speed  $s'(J_{k_0}) > s(J_{k_0})$  while all other jobs  $J_k \in \mathcal{J}_{i_0+1} \cup \dots \cup \mathcal{J}_p$  are processed at speed  $s'(J_k) \geq s(J_k)$ . We obtain a contradiction.  $\square$

In the remainder of this section we describe the change in potential arising in response to the arrival of a new job  $J_{n+1}$ . If the deadline  $d_{n+1}$  does not coincide with the deadline of a previous job  $J_i$ ,  $1 \leq i \leq n$ , we have to update the set of intervals  $I_j$ ,  $1 \leq j < |\mathcal{I}|$ . We add  $d_{n+1}$  to  $\mathcal{I}$ . If  $d_{n+1}$  is the last deadline of any of the jobs, we introduce a new interval  $[\tau_{\max}, d_{n+1})$ , where  $\tau_{\max} = \max_{1 \leq i \leq n} d_i$ ; otherwise we split the interval  $[\tau_j, \tau_{j+1})$  containing  $d_{n+1}$  into two intervals  $[\tau_j, d_{n+1})$  and  $[d_{n+1}, \tau_{j+1})$ .

In a next step we add  $J_{n+1}$  to the current schedule  $S_{OA(m)}$  assuming that the processing volume of  $J_{n+1}$  is zero. Let  $I_j$  be the interval ending at time  $d_{n+1}$ . If  $S_{OA(m)}$  has an idle processor in  $I_j$ , then we open a new job set  $\mathcal{J}_{p+1}$  with associated speed  $s_{p+1} = 0$ . Otherwise let  $s_i$  be the lowest speed used by any of the processors in  $I_j$ . We add  $J_{n+1}$  to  $\mathcal{J}_i$ . In the following we increase the processing volume of  $J_{n+1}$  from zero to its final size  $w_{n+1}$ . The increase proceeds in a series of phases such that, during a phase, the job sets  $\mathcal{J}_i$ ,  $i \geq 1$ , do not change. However, at the beginning of a phase the job set  $\mathcal{J}_{i_0}$  containing  $J_{n+1}$  might split into two sets. Moreover, at the end of a phase the set  $\mathcal{J}_{i_0}$  might merge with the next set  $\mathcal{J}_{i_0-1}$  executed at the next higher speed level. So suppose that during a phase the processing volume of  $J_{n+1}$  increases by  $\epsilon$ . By Lemma 11, the speeds of the jobs in sets  $\mathcal{J}_i$ ,  $i > i_0$ , do not change. Moreover, by Lemma 10, the speeds of the other jobs can only increase. Hence, by convexity of the power function, an optimal solution will increase the speed at the lowest possible level, which is  $s_{i_0}$ . In the current schedule of  $OA(m)$ , let  $T_{i_0}$  be the total time for which jobs of  $\mathcal{J}_{i_0}$  are processed. The speed at which jobs of  $\mathcal{J}_{i_0}$  are processed increases from  $W_{OA(m)}(i_0)/T_{i_0}$  to  $(W_{OA(m)}(i_0) + \epsilon)/T_{i_0}$ . Hence the potential change is given by

$$\begin{aligned} \Delta \Phi &= \alpha \left( \frac{W_{OA(m)}(i_0) + \epsilon}{T_{i_0}} \right) (W_{OA(m)}(i_0) + \epsilon - \alpha(W_{OPT}(i_0) + \epsilon)) \\ &\quad - \alpha \left( \frac{W_{OA(m)}(i_0)}{T_{i_0}} \right) (W_{OA(m)}(i_0) - \alpha(W_{OPT}(i_0))). \end{aligned}$$



**Algorithm AVR(m):** In each interval  $I_t$ ,  $0 \leq t < T$ , execute the following.

1.  $\mathcal{J}'_t := \mathcal{J}_t$ ;  $M$  is the set of all processors in the system;
2. **while**  $\mathcal{J}'_t \neq \emptyset$  **do**
3.     **while**  $\max_{J_i \in \mathcal{J}'_t} \delta_i > \Delta'_t/|M|$  **do**
4.         Pick a processor  $l \in M$ ; Schedule  $J_i := \arg \max_{J_i \in \mathcal{J}'_t} \delta_i$  at a speed of  $\delta_i$  on processor  $l$ ;
5.          $\mathcal{J}'_t := \mathcal{J}'_t \setminus \{J_i\}$ ;  $M := M \setminus \{l\}$ ;
6.      $s_\Delta := \Delta'_t/|M|$ ; Feasibly schedule every job  $J_i \in \mathcal{J}'_t$  at a speed of  $s_\Delta$ .

**Fig. 3.** The algorithm AVR(m).

As in [5] we can show the non-positivity of the last expression.

At the beginning of a phase  $\mathcal{J}_{i_0}$  might split into two sets  $\mathcal{J}'_{i_0}$  and  $\mathcal{J}''_{i_0}$ . Set  $\mathcal{J}'_{i_0}$  contains those jobs that, in an optimal schedule, cannot be executed at an increased speed when raising the processing volume of  $J_{n+1}$ . Set  $\mathcal{J}''_{i_0}$  contains the jobs for which the speed can be increased. This split into two sets does not change the value of the potential function. For completeness we mention that a job of  $\mathcal{J}_{i_0}$  whose execution starts, say, in interval  $I_j = [\tau_j, \tau_{j+1})$  belongs to  $\mathcal{J}'_{i_0}$  iff in the current schedule (a) no job  $J_k \in \mathcal{J}_{i_0}$  with  $d_k > \tau_j$  is processed before  $\tau_j$  or (b) any job  $J_k \in \mathcal{J}_{i_0}$  with  $d_k > \tau_j$  that is also processed before  $\tau_j$  is scheduled continuously without interruption throughout  $[\tau_j, d_k)$ . When the processing volume of  $J_{n+1}$  increases as described above, the jobs of  $\mathcal{J}'_{i_0}$  remain at their speed level  $s_{i_0}$  while the speed of the jobs in  $\mathcal{J}''_{i_0}$  increases.

At the end of a phase  $\mathcal{J}_{i_0}$  might be merged with the next set  $\mathcal{J}_{i_0-1}$ . This event occurs if speed value  $s_{i_0}$  reaches the next higher level  $s_{i_0-1}$ . We always choose  $\epsilon$  such that  $(W_{OA(m)}(i_0) + \epsilon)/T_{i_0}$  does not exceed  $s_{i_0-1}$ . Again the merge of two job set does not change the value of the potential function.

**Theorem 2.** Algorithm OA(m) is  $\alpha^\alpha$ -competitive.

### 3.2. Algorithm Average Rate

The original Average Rate algorithm for a single processor considers job densities, where the density  $\delta_i$  of a job  $J_i$  is defined as  $\delta_i = w_i/(d_i - r_i)$ . At time  $t$  the processor speed  $s_t$  is set to the sum of the densities of the jobs active at time  $t$ . Using these speeds  $s_t$ , AVR always processes the job having the earliest deadline among the active unfinished jobs.

In the following we describe our algorithm, called AVR(m), for  $m$  parallel processors. We assume without loss of generality that all release times and deadlines are integers. Moreover, we assume that the earliest release time  $\min_{1 \leq i \leq n} r_i$  is equal to 0. Let  $T = \max_{1 \leq i \leq n} d_i$  be the last deadline. For any time interval  $I_t = [t, t+1)$   $0 \leq t < T$ , let  $\mathcal{J}_t$  be the set of jobs  $J_i$  that are active in  $I_t$ , i.e. that satisfy  $I_t \subseteq [r_i, d_i)$ . Algorithm AVR(m) makes scheduling decisions for the intervals  $I_t$  over time,  $0 \leq t < T$ . In each  $I_t$ , the algorithm schedules  $\delta_i$  units of  $w_i$ , for each job  $J_i \in \mathcal{J}_t$ . The schedule is computed iteratively. In each iteration there is a set  $\mathcal{J}'_t \subseteq \mathcal{J}_t$  of jobs to be scheduled on a subset  $M$  of the processors in interval  $I_t$ . Initially  $\mathcal{J}'_t = \mathcal{J}_t$ , and  $M$  contains all the processors of the system. Furthermore, let  $\Delta_t = \sum_{J_i \in \mathcal{J}_t} \delta_i$  be the total density of jobs active at time  $t$ , and  $\Delta'_t = \sum_{J_i \in \mathcal{J}'_t} \delta_i$  be the total density of the jobs in  $\mathcal{J}'_t$ . The average load of the jobs in  $\mathcal{J}_t$  and  $\mathcal{J}'_t$  is  $\Delta_t/m$  and  $\Delta'_t/|M|$ , respectively.

In each iteration, if  $\max_{J_i \in \mathcal{J}'_t} \delta_i > \Delta'_t/|M|$ , then a job  $J_i$  of maximum density among the jobs in  $\mathcal{J}'_t$  is scheduled on a processor  $l \in M$ . Job  $J_i$  is processed during the whole interval  $I_t$  at a speed of  $s_{t,l} = \delta_i$ . Then  $\mathcal{J}'_t$  and  $M$  are updated to  $\mathcal{J}'_t \setminus \{J_i\}$  and  $M \setminus \{l\}$ , respectively. On the other hand if  $\max_{J_i \in \mathcal{J}'_t} \delta_i \leq \Delta'_t/|M|$ , the complete set  $\mathcal{J}'_t$  is scheduled at a uniform speed of  $s_\Delta = \Delta'_t/|M|$  on the processors of  $M$  as follows. First a sequential schedule  $S$  of length  $|M|$  is constructed by concatenating the execution intervals of all jobs in  $\mathcal{J}'_t$ . Again, each job in  $\mathcal{J}'_t$  is executed at a speed of  $s_\Delta$  and therefore has an execution interval of length  $\delta_i/s_\Delta$ . Then  $S$  is split among the processors of  $M$  by assigning time range  $[(\mu-1), \mu)$  of  $S$  from left to right to the  $\mu$ -th processor of  $M$ ,  $1 \leq \mu \leq |M|$ . A summary of AVR(m) is given in Fig. 3.

It is easy to see that AVR(m) computes feasible schedules. In any interval  $I_t$  where a job  $J_i$  is active,  $\delta_i$  processing units of  $J_i$  are finished. Summing over all intervals where  $J_i$  is active, we obtain that exactly  $\delta_i(d_i - r_i) = w_i$  processing units, i.e. the total work of  $J_i$ , are completed. Furthermore, at no point in time is a job scheduled simultaneously on two different processors: In any interval  $I_t$ , a job  $J_i \in \mathcal{J}_t$  is either scheduled alone on just one processor or it is scheduled at a speed of  $s_\Delta$ . In the second case, since  $J_i$  has a density of  $\delta_i \leq s_\Delta$  its execution interval in  $S$  has a length of at most one time unit. Therefore, even if  $J_i$  is split among two processors  $\mu$  and  $\mu+1$ ,  $1 \leq \mu < |M|$ , it is processed at the end of processor  $\mu$  and at the beginning of  $\mu+1$  so that these two execution intervals do not overlap (recall that  $|I_t| = 1$ ).

**Theorem 3.** AVR(m) achieves a competitive ratio of  $(2\alpha)^\alpha/2 + 1$ .

**Proof.** Let  $\sigma = J_1, \dots, J_n$  be an arbitrary job sequence. The energy incurred by AVR(m) is

$$E_{AVR(m)}(\sigma) = \sum_{t=0}^{T-1} \sum_{l=1}^m (s_{t,l})^\alpha |I_t|.$$

The length  $|I_t|$  of each interval  $I_t$  is equal to 1 time unit, but we need to incorporate  $|I_t|$  in the above expression in order to properly evaluate the consumed energy.

For each  $I_t$ , we partition the  $m$  processors into two sets depending on whether or not the actual speeds are higher than  $\Delta_t/m$ . More precisely, let  $M_{t,1}$  be the set of all  $l$ ,  $1 \leq l \leq m$ , such that the speed of processor  $l$  is  $s_{t,l} \leq \frac{\Delta_t}{m}$ . Let  $M_{t,2} = \{1, \dots, m\} \setminus M_{t,1}$ . We have

$$\begin{aligned} E_{AVR(m)}(\sigma) &\leq \sum_{t=0}^{T-1} \sum_{l \in M_{t,1}} \left( \frac{\Delta_t}{m} \right)^\alpha |I_t| + \sum_{t=0}^{T-1} \sum_{l \in M_{t,2}} (s_{t,l})^\alpha |I_t| \\ &\leq \sum_{t=0}^{T-1} m^{1-\alpha} (\Delta_t)^\alpha |I_t| + \sum_{t=0}^{T-1} \sum_{l \in M_{t,2}} (s_{t,l})^\alpha |I_t|. \end{aligned}$$

The last inequality holds because  $|M_{t,1}| \leq m$ . Next, for any fixed  $I_t$ , consider the processors of  $M_{t,2}$ . By the definition of  $AVR(m)$  any such processor is assigned exactly one job. Hence, for any processor  $l$  with  $l \in M_{t,2}$ , the speed  $s_{t,l}$  is the density of the job assigned to it. Hence  $\sum_{l \in M_{t,2}} (s_{t,l})^\alpha \leq \sum_{J_i \in \mathcal{J}_t} (\delta_i)^\alpha$  and we obtain

$$E_{AVR(m)}(\sigma) \leq m^{1-\alpha} \sum_{t=0}^{T-1} (\Delta_t)^\alpha |I_t| + \sum_{t=0}^{T-1} \sum_{J_i \in \mathcal{J}_t} (\delta_i)^\alpha |I_t|. \quad (9)$$

In the above expression  $\sum_{t=0}^{T-1} \sum_{J_i \in \mathcal{J}_t} (\delta_i)^\alpha |I_t|$  each job contributes an energy of  $(\delta_i)^\alpha (d_i - r_i)$  because  $J_i \in \mathcal{J}_t$  for any  $t \in [r_i, d_i)$ . This amount is the minimum energy required to process  $J_i$  if no other jobs were present. Hence  $\sum_{t=0}^{T-1} \sum_{J_i \in \mathcal{J}_t} (\delta_i)^\alpha |I_t| \leq E_{OPT}(\sigma)$ , where  $E_{OPT}(\sigma)$  is the energy consumption of an optimal schedule for  $\sigma$  on the  $m$  processors.

As for the first term in (9),  $\sum_{t=0}^{T-1} (\Delta_t)^\alpha |I_t|$  is the energy consumed by the single processor algorithm  $AVR$  when executing  $\sigma$  on one processor. Recall that  $AVR$  sets the speed at time  $t$  to the accumulated density of jobs active at time  $t$ . Since  $AVR$  achieves a competitive ratio of  $(2\alpha)^\alpha/2$ , see [15], we obtain  $\sum_{t=0}^{T-1} (\Delta_t)^\alpha |I_t| \leq \frac{1}{2} (2\alpha)^\alpha E_{OPT}^1(\sigma)$ , where  $E_{OPT}^1(\sigma)$  is the energy of an optimal single processor schedule for  $\sigma$ . Thus

$$E_{AVR(m)}(\sigma) \leq m^{1-\alpha} \frac{1}{2} (2\alpha)^\alpha E_{OPT}^1(\sigma) + E_{OPT}(\sigma).$$

To establish the desired competitive ratio, we finally prove  $m^{1-\alpha} E_{OPT}^1(\sigma) \leq E_{OPT}(\sigma)$ .

Consider an optimal schedule  $S_{OPT}$  for  $\sigma$  on  $m$  processors and let  $s_{t,l}^o$  be the speed on processor  $l$  in  $I_t$ , where  $1 \leq l \leq m$  and  $0 \leq t < T - 1$ . Let  $s_t^o = \sum_{l=1}^m s_{t,l}^o$  be the sum of the speeds in  $I_t$ . By the convexity of the power function

$$E_{OPT}(\sigma) = \sum_{t=0}^{T-1} \sum_{l=1}^m (s_{t,l}^o)^\alpha |I_t| \geq \sum_{t=0}^{T-1} m (s_t^o/m)^\alpha |I_t| = m^{1-\alpha} \sum_{t=0}^{T-1} (s_t^o)^\alpha |I_t|. \quad (10)$$

Finally, consider the single processor schedule  $S^1$  that, in any interval  $I_t$ , uses speed  $s_t^o$  and accomplishes the same work as  $S_{OPT}$  in  $I_t$ . More precisely, if  $S_{OPT}$  finishes  $w_{t,i}$  processing units of  $J_i$  in  $I_t$ , where  $0 \leq w_{t,i} \leq w_i$ , then  $S^1$  processes the same amount of  $J_i$  in  $I_t$ . Using speed  $s_t^o$ , this is indeed possible. Hence  $S^1$  is a feasible schedule whose energy consumption is  $\sum_{t=0}^{T-1} (s_t^o)^\alpha |I_t|$ , and this expression is at least  $E_{OPT}^1(\sigma)$ . Combining with (10) we obtain  $E_{OPT}(\sigma) \geq m^{1-\alpha} \sum_{t=0}^{T-1} (s_t^o)^\alpha |I_t| \geq m^{1-\alpha} E_{OPT}^1(\sigma)$ .  $\square$

#### 4. Conclusion and open problems

In this paper we have investigated speed scaling in multi-processor environments. We considered a classical scheduling problem where jobs have associated deadlines and assumed that job migration is allowed. We have developed a combinatorial algorithm for computing optimal offline schedules and extended the single processor algorithms *Optimal Available* and *Average Rate*. Bansal et al. [5] gave an online algorithm for a single processor that, for large  $\alpha$ , achieves a smaller competitiveness than *Optimal Available*. An open problem is if this strategy can also be extended to multi-processor systems.

Another working direction is to devise and analyze online algorithms for general convex power functions. Even for a single processor, no competitive strategy is known. Moreover, Irani et al. [9] investigated a problem setting where a single variable-speed processor is equipped with an additional sleep state into which the processor can be transitioned when idle. The framework is interesting because modern microprocessors, even at speed zero, consume a significant amount of static energy. It would be worthwhile to investigate combined speed scaling and power-down mechanisms in multi-processor environments.

## References

- [1] S. Albers, F. Müller, S. Schmelzer, Speed scaling on parallel processors, *Algorithmica* 68 (2) (2014) 404–425.
- [2] N. Bansal, D.P. Bunde, H.-L. Chan, K. Pruhs, Average rate speed scaling, *Algorithmica* 60 (4) (2011) 877–889.
- [3] N. Bansal, H.-L. Chan, T.-W. Lam, L.-K. Lee, Scheduling for speed bounded processors, in: *Proc. 35th International Colloquium on Automata, Languages and Programming, ICALP*, in: LNCS, vol. 5125, Springer, 2008, pp. 409–420.
- [4] N. Bansal, H.-L. Chan, K. Pruhs, D. Katz, Improved bounds for speed scaling in devices obeying the cube-root rule, *Theory Comput.* 8 (1) (2012) 209–229.
- [5] N. Bansal, T. Kimbrel, K. Pruhs, Speed scaling to manage energy and temperature, *J. ACM* 54 (2007).
- [6] B.D. Bingham, M.R. Greenstreet, Energy optimal scheduling on multiprocessors with migration, in: *Proc. IEEE International Symp. on Parallel and Distributed Processing with Applications*, 2008, pp. 143–152.
- [7] H.-L. Chan, W.-T. Chan, T.-W. Lam, L.-K. Lee, K.-S. Mak, P.W.H. Wong, Optimizing throughput and energy in online deadline scheduling, *ACM Trans. Algorithms* 6 (1) (2009).
- [8] G. Greiner, T. Nonner, A. Souza, The bell is ringing in speed-scaled multiprocessor scheduling, *Theory Comput. Syst.* 54 (1) (2014) 24–44.
- [9] S. Irani, S.K. Shukla, R. Gupta, Algorithms for power savings, *ACM Trans. Algorithms* 3 (2007).
- [10] T.-W. Lam, L.-K. Lee, I.K.-K. To, P.W.H. Wong, Energy efficient deadline scheduling in two processor systems, in: *Proc. 18th International Symposium on Algorithms and Computation*, in: LNCS, vol. 4835, Springer, 2007, pp. 476–487.
- [11] M. Li, B.J. Liu, F.F. Yao, Min-energy voltage allocation for tree-structured tasks, *J. Comb. Optim.* 11 (2006) 305–319.
- [12] M. Li, A.C. Yao, F.F. Yao, Discrete and continuous min-energy schedules for variable voltage processors, *Proc. Natl. Acad. Sci. USA* 103 (2006) 3983–3987.
- [13] M. Li, F.F. Yao, An efficient algorithm for computing optimal discrete voltage schedules, *SIAM J. Comput.* 35 (2005) 658–671.
- [14] D.D. Sleator, R.E. Tarjan, Amortized efficiency of list update and paging rules, *Commun. ACM* 28 (1985) 202–208.
- [15] F.F. Yao, A.J. Demers, S. Shenker, A scheduling model for reduced CPU energy, in: *Proc. 36th IEEE Symposium on Foundations of Computer Science*, 1995, pp. 374–382.